

General Knowledge

- Structure: End points ↔ Comm Links ↔ Switch ↔ End points

End-points (Network Edge)	PC, Servers, Phones, ...
Comm Links (Access Net / Phys Media)	Copper, Fiber, WiFi, 5G, ...
Packet Switches (Network Core)	Router, Switches, ...

High-level Structure:

Network of Networks	Interconnected ISP
Network Protocols	Rules that govern how data exchange
Internet Standards	Ensure diff. HW & SW work together

Access Network

Digital Subscriber Line (DSL)	Use existing telephone lines. Direct link to telecom company's central office (dedicated access)
Cable Networks	Fiber coax cable: homes ↔ ISP (shared access)
Ethernet	Fiber optic cables. Used in institutions / company
Wireless Access Network	Wireless LAN : short range. E.g., WiFi Widearea Wireless Access E.g., Cellular

Physical Media

Coaxial Cable	(Guided Media): 2 concentric copper conductors <ul style="list-style-type: none">Bidirectional, broadband (multiple channel)
---------------	---

Fiber Optical Cable	(Guided Media): Glass fiber, light pulses. High Speed <ul style="list-style-type: none">Low err. rate. Immune electromagnetic noise
---------------------	--

Radio Signal	(Unguided Media): electromagnetic signals <ul style="list-style-type: none">Stability: reflection / obstruction / interferenceWireless, Bidirectional
--------------	---

Network Core

Packet Switching:


- Host break applicaiton-layer message into *packets*.
- Host send packets to edge router
- Router forward one to the next (**wait for entire packet arrive**)

Key Functions

- Routing: determine source-destination route
- Forwarding: move packet to next router (input → output)

Switching Modes

Frequency Division Multiplexing (FDM) & Time Division ... (TDM)



Transmission Loss & Delay

Delay at 1 Node: $d_{\text{nodal}} = d_{\text{proc}} + d_{\text{queue}} + d_{\text{trans}} + d_{\text{prop}}$

- d_{proc} **processing delay**: check bit error, determine output link
- d_{queue} **queueing delay**: time waiting at output link (congestion)
- d_{trans} **transmission delay**: $d_{\text{trans}} = \sum_i \frac{\text{packet size (bits)}}{\text{transmission rate at hop } i \text{ (bits/sec)}}$
- d_{prop} **propagation delay**: $d_{\text{prop}} = \sum_i \frac{\text{length of link } i}{\text{propagation speed in medium}}$

More on Queueing Delay:

L : packet length (bits); R : link bandwidth (bps); a : average arrival rate

Avg service time	$T = \frac{L}{R}$	Notice,
Network Intensity	$\rho = \frac{L \cdot a}{R} = T a$	$\rho \rightarrow 0, d_{\text{queue}} \rightarrow 0,$
Avg queueing delay	$d_{\text{queue}} = \frac{T \rho}{1 - \rho} = \frac{L \rho}{R(1 - \rho)}$	$\rho \rightarrow 1, d_{\text{queue}} \rightarrow \infty$

Packet Loss

When $L a > R$: Packet queue at router. Packet lost if memory full.

Throughput

Rate (**bits/time unit**) which bits transferred between sender/receiver

- instantaneous: rate at given point in time
- average: rate over longer period of time

Theoretical Upperbound: **max-flow min-cut** of the network

Protocol Layers & Service Model

TCP/IP

Application	Application-specific functionality (Data) FTP, HTTP
Transport	Process ↔ process w/ports (Segments) TCP, UDP
Network	Host ↔ host w/IP (Packets) IP, routing protocols.
Link	Node ↔ node w/MAC (Frame) Ethernet, WiFi, PPP
Physical	Binary Bits "in" physical medium

ISO/OSI Reference Model

Seven Layers Application → Presentation → Session → Transport → Network → Data Link → Physical

Two more than TCP/IP:

- Presentation: data format, encryption, compression
- Session: synchronization, checkpointing, recovery of data


Application Layer

- Non-persistent HTTP:
 - Each object request → new TCP connection

- 2 RTT to fetch every object


- Persistent HTTP
 - Server leaves TCP connection open after init
 - Object sent over same TCP connection
 - Minimum of **1 RTT** to fetch all objects

HTTP Request



- cr: carriage return; lf: line feed
- Method types:
 - HTTP/1.0: GET, POST, HEAD
 - HTTP/1.1: GET, POST, HEAD, PUT (uploads file in entity body to path specified in URL field) DELETE (deletes file specified in URL field)
 - POST method: data in request message body. High security.
 - URL (or GET) method: data in URL field. Low security.

HTTP Response



- Status code:
 - 200: OK, 301: Moved Permanently, 400: Bad Request, 404: Not Found, 505: HTTP Version Not Supported

Caching

- Total delay = Internet delay + access delay + LAN delay
- HTTP request: If-modified-since: <date>
- HTTP response: 304 Not Modified or 200 OK

Email

- Send: SMTP (Simple Mail Transfer Protocol)
- Recv: POP3 (Post Office Prot) / IMAP (Internet Mail Access Prot)
- User Agent → Mail Server A → Mail Server B → User Agent
- CRLF. CRLF to end message

POP3 (Post Office Protocol): Deletes message, stateless

- Authorization phase:
 - Client command to server: user <username>, pass <password>
 - Server response: +OK, -ERR
- Transaction phase:
 - list: list message numbers
 - retr <msg_number>: retrieve message by number
 - dele <msg_number>: delete message by number
 - quit: end session

IMAP (Internet Mail Access Protocol): Keep state

- Keep all messages on server
- Allows user to organize messages in folders

DNS (Domain Name System)

- Root Name Servers: 13 root servers. return IP of TLD server
- TLD Name Servers: return IP of authoritative name server
 - Responsible for com, org, net, edu, ... uk, fr, ca, ...
- Authoritative Name Servers: return IP of host
 - Usually organization's own DNS server (e.g., Cloudflare)
- Local DNS: each ISP has one. not part of hierarchy.
 - If have cache, return (might outdated). Else, forward into hierarchy.

Resource Records (RR): format (name, value, type, ttl)

- A: hostname → IP. name = hostname, value = IP
- NS: name → auth DNS. name = domain, value = auth DNS hostname
- MX: mail server for domain. name = domain, value = server name
- CNAME: canonical name. name = alias, value = real name

Transport Layer

- Multiplexing: assign app data to different ports, add these port headers to outgoing packets. Let multiple apps to use same network.
- Demultiplexing: based on incoming packet port numbers, direct data to correct app.

Reliable Data Transfer (RDT)

- RDT 1.0: do nothing, assume no error
- RDT 2.0: solve packet corruption.
 - Adds checksum to detect error.
 - Receiver send ACK/NAK to acknowledge sender.
- RDT 2.1: solve ACK/NAK packet corruption
 - Important: assumes stop-and-wait protocol (1 packet at a time)
 - Sender Rule:
 - Send packet 0, wait for ACK
 - If NAK or ACK/NAK corrupt, resend packet 0
 - If ACK, send packet 1, wait for ACK
 - Receiver Rule:
 - If packet corrupt, send NAK
 - If packet 0 received, when prev is 1, send ACK, flip 0/1

- If packet 1 received, when prev is 0, send ACK, flip 0/1
- If packet 0 received two times, delete duplicate, send ACK again

- RDT 2.2: NAK-free protocol - Send ACK 0/1 instead of NAK
 - Receiver ACK just the last correct packet if new packet corrupt
 - Sender retransmit packet if duplicate ACK received
- RDT 3.0: Add timeout to sender (Solves ACK packet loss)

Performance of RDT 3.0:

- Transmission Time = $D_{\text{trans}} = \frac{\text{Length}}{\text{Rate}} + \text{RTT}$
- Utilization (sender) = $U_{\text{sender}} = \frac{\text{RTT} + L/R}{\text{RTT} + L/R}$

Pipelining: Utilization (sender): $U_{\text{sender}} = \frac{N \times L/R}{\text{RTT} + L/R}$

- where N is the number of packet transmitted at once (parallelization)

Go-Back-N (GBN):

- Key Idea: ensure **in-order delivery**.
 - Receiver discard out-of-order pkt, only ACK packet $i - 1$
 - Cumulative ACK sender get any ACK i , mark all pkts $\leq i$ ACK
- Sender:
 - Send N (window size) packets, wait for ACK
 - Receive any valid ACK i , move window to $[i + 1, i + N]$
 - If 1st pkt in window timeout, resend entire window
- Receiver:
 - Only accept next in-order pkt (window size = 1, conceptually)
 - Only accept pkt i if and only if i is the next seq#
 - If got out-of-order pkt $j > i$, discard and send ACK $i - 1$
 - If got in-order pkt i , send ACK i


Prevent Corruption: $N \leq K - 1$, where K is the #seq range

Selective Repeat (SR):

Prevent Corruption: $N \leq \lfloor K/2 \rfloor$, where K is the #seq range

- Sender:
 - Send N (window size) packets, wait for ACK
 - Track timeout of **every** single packet. Resend whichever timeouts
- Receiver:
 - ACK (not cumulative) every single valid pkt (even out-of-order)
 - Buffer out-of-order packets within its window (window size = N)
 - Send buffered pkt to application layer once prior pkts arrive.

TCP (Transmission Control Protocol)



- Sequence Num:
 - Byte number of 1st byte in the current segment
 - NOT the same as #seq in GBN and SR
 - Example: SEQ=1000 data length=500 bytes (Sent bytes 1000-1499)
- Acknowledgement Num:
 - Sequence number of next byte expected from the other side
 - Example (Con't): RECV ACK=1500 (i.e., next byte expected is 1500)
- 3-way-handshake w/ SYN and ACK flags
 - A → B: SYN, seq = x A initiate connection
 - B → A: SYN-ACK, seq = y, ack = x+1 B acknowledge A's SYN
 - A → B: ACK, ack = y+1. A acknowledge B's SYN
 - A and B maintain randomly initialized unique seq #

TCP Timeout Control

EstimateRTT_n = (1 - α) · EstimateRTT_{n-1} + α · SampleRTT_n
DevRTT_n = (1 - β) · DevRTT_{n-1} + β · |SampleRTT_n - EstimateRTT_n|
TimeoutInterval_n = EstimateRTT_n + 4 · DevRTT_n

- where α = 0.125, β = 0.25 typically, EMA
- Timeout = avg RTT + 4 * deviation (for safety margin)
 - too short: premature timeout, unnecessary retransmission
 - too long: slow reaction to segment loss

TCP Receiver Actions:

- Recv in-order seg, previous received all already ACK'ed
 - Action: Wait 500ms for new pkt, if no, send ACK of current pkt.
- Recv in-order seg, previous received pkt not yet ACK'ed
 - Action: immediately send **single** accumulative ACK (all previous)
- Recv out-of-order seg $j > i$ (Gap Detected)
 - Action: immediately send **duplicate** ACK $i - 1$, with seq# = i
- Recv seg that fills **lower-end** gap caused in the above case.
 - Action: immediately send ACK.
- If sender recv 3 duplicate ACK of same data, resend un-ACK'ed segment with smallest seq # immediately.
- Receiver "advertise" free buffer space via **rwnd** value,
 - sender limit in-flight data in response to this value

TCP Congestion Control

- Additive Increase, Multiplicative Decrease (AIMD):
 - Increase cwnd by 1 MSS (Maximum Segment Size) **every** RTT.
 - Cut cwnd by half when loss detected.
 - rate ≈ cwnd ÷ RTT (bytes/sec)
- TCP Reno:
 - timeout loss: cwnd = 1 MSS, exp grow to ssthresh, then lin grow

- 3 duplicate ACK loss: ssthresh and cwnd ÷ 2, then linearly grow
- **TCP Tahoe**: always set cwnd to 1 MSS when loss detected
- **avg TCP throughput**: $\frac{3}{4} \cdot \frac{\text{Window Size}}{\text{RTT}}$ bytes/sec

Network Layer

- **forwarding (data plane)**: move packets from a router's input to output
- **routing (control plane)**: determine route from source to destination
 - *Traditional routing algorithms*: implemented in routers (local decision)
 - *Software-defined networking (SDN)*: implemented in (remote) servers

Router

Input Port (3 components):


1. Line termination (physical layer), bit-level reception
2. link layer protocol (data link layer, e.g., Ethernet)
3. Forwarding:
 - Use header field values → determine output port
 - If datagram arrive fast, queue at **input port buffer**

Input Port Queuing

- **Head-of-the-line (HOL) blocking**: in same input port, packet in front of line blocking packets behind (e.g., when head packet's output port is busy)

Switching Fabric (3 types of methods):

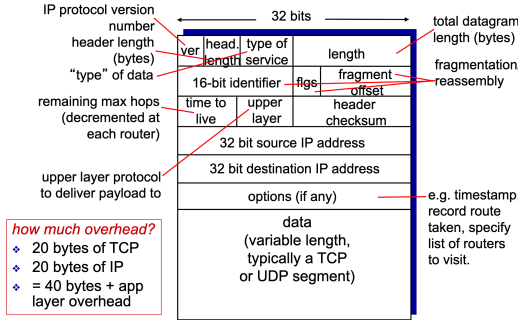
1. **Memory**: CPU direct control. packet copied to RAM, speed bottlenecked.
2. **Bus**: datagram move thru. shared bus. Limited by bus bandwidth.
3. **Crossbar**: high parallelism, n^2 crosspoints.



Output Port (3 components):

1. Datagram buffer. Required when datagram arrives > transmission rate
2. Link layer protocol (send)
3. Line termination

Internet Protocol (IP)



IPv4 Address: 32-bits, 2 levels, 5 classes (A, B, C, D, E)

- A: begin w/ (0)₂ 0-127, 8/24bit, $\frac{1}{2}$ addr space, mask: 255.0.0.0
- B: begin w/ (10)₂ 128-191, 16/16bit, $\frac{1}{4}$ addr space, mask: 255.255.0.0
- C: begin w/ (110)₂ 192-223, 24/8bit, $\frac{1}{8}$ addr space, mask: 255.255.255.0
- D: begin w/ (1110)₂ 224-239, multicast, $\frac{1}{16}$ addr space
- E: begin w/ (1111)₂ 240-255, reserved (future proof), $\frac{1}{16}$ addr space

Subnets: devices / interfaces with same netid (network part of IP address), physically reach each other without router.

Classless Inter-Domain Routing (CIDR):

- subnet part (netid) is arbitrary length
- format: **a.b.c.d/x**, where **x** is the # bits in subnet portion of address

IP Fragmentation:

Purpose:

- Network links have **MTU (Max Transmission Unit)**
- Need to fragment datagrams, then reassemble at destination.

Header:

- length field: total length of datagram (header + data)
- offset field: data length ÷ 8
- fragflag: 0 for last fragment, 1 for not last

Destination-based Forwarding:

- Longest prefix matching: forward table contains list of prefixes, find longest prefix match

DHCP (Dynamic Host Configuration Protocol):

Purpose: let host dynamically get IP addr. from network server when join

- Host: broadcast **DHCP discover** msg to find DHCP server
- DHCP server: respond **DHCP offer** msg
- Host: requests IP address via **DHCP request** msg
- DHCP server: send **DHCP ack** msg, which returns:
 - allocated IP, first-hop router addr, network mask, name and IP of DNS

ICANN (Internet Corporation for Assigned Names and Numbers):

- Assign blocks of IP addresses to the RIRs (Regional Internet Registries); RIRs assign to ISPs or large organizations

Routing Protocols

Bellman-Ford Algorithm:

- Same as Dijkstra's algorithm: $d(x) = \min(d(x), d(y) + l(y, x))$
- Except, don't choose min node each iteration.

- for $n - 1$ iterations ($n = \text{number of nodes}$),
 - **for each m edge (y, x) :**
 - $d(x) = \min(d(x), d(y) + l(y, x))$ **if** $y \rightarrow x$
 - $d(y) = \min(d(y), d(x) + l(x, y))$ **if** $x \rightarrow y$

Distance Vector Algorithm:

- Same as Bellman-Ford, except asynchronous, distributed, iterative
- Maintains a 多源最短路矩阵
- Each node:
 1. **wait** for change in local link cost, or msg from neighbor
 2. **re-compute** distance vector estimate using Bellman-Ford
 3. if DV to **any** dest changed, notify **all** neighbors
- **Characteristics**:

• **Good news travel fast** (adj edge cost ↓): node cost unchange / improve. Advertise to all neighbors. Update ripple out one hop at a time (fast).

• **Bad news travel slow** (adj edge cost ↑): **A** drop routes that use the bad edge, seek alternative from neighbor **B**. But neighbor **B** might route through **A**, its shorter only because it is outdated. End up "bouncing" distance updates back and forth (slow, *count-to-infinity problem*).

- **Solution**: if **B** routes through **A**, set **B**'s distance to ∞ .

Link-State Routing (Dijkstra) vs. Distance Vector (Bellman-Ford)

	LS (Dijkstra)	DV (Bellman-Ford)
Message Complexity	$O(nm)$	exchange msg between adj nodes only
Convergence Speed	$O(n^2)$	vary. loops / count-to-infinity

AS: Autonomous System, group of routers under same admin control

Intra-AS Routing:

- **Interior Gateway Protocols (IGPs)**:
 - **RIP (Routing Information Protocol)**:
 - distance metric: **number of hops** (max 15 hops)
 - update interval: DV advertise every 30 seconds
 - routing table managed by application-level process route-d (daemon)
 - **Open Shortest Path First (OSPF)**:
 - LS-based, use Dijkstra's algorithm
 - OSPF message directly over IP (**not** TCP/UDP). Use protocol field 89
 - All OSPF messages authenticated; Multiple same-cost paths allowed
 - Each link, multiple cost metric for different services
 - Supports multicast (MOSPF) and unicast
 - **Hierarchical OSPF** in large domains

Inter-AS Routing:

- Forwarding table configured by both intra- and inter-AS routing algorithm
- **Hot Potato Routing**:
 - When destination reachable via multiple AS, greedily choose gateway w/ least cost. Don't care about cost of inter-AS routes.
- **BGP** (Border Gateway Protocol):
 - Types:
 - **eBGP**: obtain subnet reachability from neighboring AS
 - **iBGP**: propagate reachability to AS-internal routers
 - BGP Messages:
 - OPEN: open semi-persistent TCP connection to remote BGP peers
 - UPDATE: advertises new paths (or withdraw old)
 - KEEPALIVE: keep connection alive in absence of UPDATE, and ack OPEN
 - NOTIFICATION: report errors in previous msg; also for closing conn.
 - BGP Route Selection:
 1. Shortest AS-path
 2. Closest next-hop router
 3. local preference value attribute: policy-based selection
 4. Additional criterias

Link Layer

Parity Checking:

- **Single Bit Parity**: **detect** single bit error *no errors*
- **Two-Dimensional Parity**: **detect** and **correct** single bit error

Cyclic Redundancy Check (CRC):

1. Choose divisor **G** (**must** be $r + 1$ bits long)
2. Left shift data bits **D** by r bits $\Rightarrow D \cdot 2^r$
3. Obtain remainder: $R = (D \cdot 2^r) \bmod G$ (r bits) $R = (D \cdot 2^r) \bmod G$
4. Form final code: concatenate **D** with **R** (equivalently, $D \cdot 2^r \text{ xor } R$)

• **IMPORTANT**: both addition in " nG " and subtraction in " $\div G^n$ is **XOR**

• **Intuition**:

- Normally: $A * B = (A_0 \cdot 2^0 \cdot B) + (A_1 \cdot 2^1 \cdot B) + \dots + (A_n \cdot 2^n \cdot B)$
- But in CRC: $A * B = (A_0 \cdot 2^0 \cdot B) \text{ xor } \dots \text{ xor } (A_n \cdot 2^n \cdot B)$
- Normally in long division, subtract each $(A_i \cdot 2^i \cdot B)$ from dividend
- But in CRC, we xor each $(A_i \cdot 2^i \cdot B)$ from dividend
- Intuitively, both + and - replaced with **xor**.
- Crucially, this works because **xor is its own inverse**

Multiple Access Protocols:

- Ideally goals: **1)** When 1 node, send at rate **R** (channel max rate) **2)** When **M** nodes, each send at average rate **R/M** **3)** Fully decentralized: no central controller / clock for synchronization **4)** Simple
- **MAC (medium access control) protocols**:
 1. **Channel Partitioning MAC protocols**:
 1. TDMA (Time Division Multiple Access):

- Channel time divided into fixed length slots
 - Each station gets a slot; unused slots go idle
2. **FDMA** (Frequency Division Multiple Access):
 - Channel spectrum divided into frequency bands
 - Each station assign a fixed frequency. Unused freq bands go idle
 2. **Random Access MAC protocols**:
 1. **Slotted ALOHA**:
 - Each node: attempt to transmit as soon as possible
 - When no collision: transmit at full channel rate
 - When collision: no node can send (slot wasted)
 - each node retry in subsequent slot with prob. **p** until success
 - **Assumptions**: all frame same size; time divided into equal-sized slots; nodes are synchronized, send only at time slot boundary;
 - **Pros**: single node full rate; highly decentralized; simple
 - **Cons**: collision waste slots; clock synchronization needed
 - **Low efficiency**: $\lim_{N \rightarrow \infty} Np(1-p)^{N-1} = \frac{1}{e} \approx 37\%$, for **N** nodes each transmit in a slot with probability **p**
 2. **Pure (unslothed) ALOHA**:
 - Each node send ASAP (no slot, no boundary, thus no sync needed)
 - Collision when ≥ 2 nodes overlap. **Worse efficiency**:
 - $\lim_{N \rightarrow \infty} Np(1-p)^{2(N-1)} = \frac{1}{2}e \approx 18\%$
 3. **CSMA** (Carrier Sense Multiple Access):
 - Listen before transmit. Send if idle; defer if busy.
 - Collision still happen (two nodes think channel is idle)
 4. **CSMA/CD** (Carrier Sense Multiple Access with Collision Detection):
 - Added collision detection. Better performance than ALOHA
 3. **"Taking turns" MAC protocols**:
 1. Polling: master node polls one by one from "slave" devices
 - Polling overhead, latency, single point of failure
 2. Token ring: token passed around ring. Only token-holder can send
 - Token overhead, latency, single point of failure (token)

Local Area Networks (LANs)

MAC Address (aka LAN Address)

- used "locally" to send frame between interfaces (same network)
- 48 bit address, burned in NIC ROM, hex (e.g., 1A-2F-BB-76-09-AD)
- Alloc. administered by IEEE, manufacturer obtain portion of addr space
- **ARP** (addr. resolution protocol): ip → mac; **RARP** (reverse ARP): mac → ip
 - each node maintain a ARP cache (ARP table: <IP, MAC>), with TTL
 - A *broadcast* ARP request containing B's IP address
 - B receives ARP packet and *unicast* reply packet to A

• **Addressing**:

1. datagram created with IP src/dst, unchanged throughout transmission
2. At each node: lookup next-hop IP → lookup MAC in ARP table → add self MAC and next-hop MAC to frame header as MAC src and MAC dst

Ethernet

- preamble: **7 bytes** of (10101010)₂ + **1 byte** of (10101011)₂
 - (10101010)₂ creates fixed wave pattern, let receiver sync w/ sender freq.
 - (10101011)₂ known as *Start Frame Delimiter* (SFD), tells receiver to start
- dest address and src address: **6 bytes** each (contains MAC address)
 - If interface receives a frame w/ dest addr not matching its own, discard
- **type**: **2 bytes** upper layer protocol (e.g., mostly IP)
- **data (payload)**: **46-1500 bytes**
- **CRC**: **4 bytes**. If error detected, frame dropped

Ethernet Characteristics:

- Connectionless & unreliable: no handshake, no ack. It's upper layer's job
- Uses **Unslotted CSMA/CD w/ binary backoff**
 1. NIC receives datagram from network layer ⇒ create frame
 2. If sense channel idle, transmit frame. Else, wait until idle
 3. If transmit w/o collision, done; Otherwise, abort.
 4. After abort, **exponential backoff**: for m -th collision, wait $K \times 512$ bit time (K sample from $\{0, 1, \dots, 2^{m-1}\}$). Then return to **step 2** above

Switches

1. **Hub**: bits come in on one port ⇒ broadcast to **all** other ports at same rate
 - All connecting nodes can collide w/ one another
 - No frame buffering; no CSMA/CD at hub: host NICs detect collision
2. **Switch**: examine incoming frame's dest address ⇒ selective forward
 - Buffer incoming frames at switch. CSMA/CD before forwarding.
 - Transparent (invisible to host); Plug-and-play (no config needed)
 - **Switch table**:
 - <MAC, Port, TTL>: records which port reach the host with the MAC
 - Learning: when recv frame, record sender location in switch table

Switch Update & Forwarding

1. **Assume**: frame **F** received at switch **S** from port **p**
2. **Given**: switch table at **S**.table
3. src, dst ← unpack(**F**) // src and dst MAC addr in **F**
4. **S**.table[src] = **p**
5. **if** dst in **S**.switch_table **then**:
 6. **if** **S**.table[dst] == **S**.table[src] **then** drop **F**
 7. **else** forward **F** to interface at **S**.table[dst]
 8. **else** flood // forward **F** to all interface, except arriving interface