

Introduction

OS Components

Level 1: Kernel

- Process scheduling, memory allocation, syscall ...

Level 2: Syscalls

- Interface between user and applications and kernel
- File I/O: `open()`, `write()` ...
- Process Control: `exec()`, `exit()`, `fork()` ...
- Communication: `ssocket()`, `send()`, `recv()` ...

Level 3: Shell & System Programs

- Shell: CLI. Translates commands → syscalls
- System Programs: Apps built on syscalls. (compilers, fork format etc)



Process Management

Process Management: Create, delete, suspend, resume, schedule

- **Multiprogramming:** Multiple programs loaded in RAM at the same time (switch when I/O bound)
- **Multitasking:** Rapidly between processes on time slice (illusion of simultaneous execution)

Memory Management: Register < Cache < RAM < 2nd Storage

- Tasks: Allocate / Deallocate, Protect unauthorized access, Optimization (paging, swapping, caching)

Storage (File) Management: Abstract storage into files / directories

- Task: Create / delete / open files. Enforce permissions (rwx). Backup, fragmentation management.

Device Management I/O Subsystem & Mass-Storage Management

- **IO subsystem:** Buffering (temporary storage for data in transit), Caching (store frequently used data in RAM), Spooling (store data for later use), Device Drivers (interface between OS and device)
- **Mass-Storage:** Free-space management, Storage allocation, Disk scheduling (minimize head movements)

User Command Interface: CLI, GUI. Translate user input → syscall

Interrupts & System Calls

Interrupt Processing

Basic Idea: Major event → Signal to the CPU → Seize the CPU **Priority:**

- **Maskable Interrupt:** low priority. Ignored / handled later
- **Non-Maskable Interrupt (NMI):** high priority. Handled immediately
 - Higher priority NMI, can interrupt lower priority NMI
 - Otherwise, low priority NMI waits for prev NMI to finish

Procedure:

1. Non-executable interrupt occurs
2. Save current CPU state (PC, Register) to stack
3. Look up an **interrupt vector table (IVT)** to find appropriate ISR
4. [User → Kernel Mode]: Jump to **Interrupt Service Routine (ISR)**
5. [**In Kernel Mode**]: Execute ISR
6. Restore CPU state from stack
7. Re-enable interrupts

Interrupt Types

- **Windows:** Ctrl+Alt+Del → Secure Attention Sequence (SAS) → Hardware Interrupt. Major, unique system-level interrupt.
- **Unix:** Ctrl+C → Software Interrupt (User Trap). General purpose software termination
- Hardware Interrupts are reliable. BUT, impossible to catch
- Software Interrupts let program handle interrupt (cleanup resource)




Figure 1: Dual Mode Operation: distinguished by *mode bit*

[2] I/O Interrupt




Figure 2: CPU interrupt needed 1 tick after transfer complete

Must know calculations:

- Sync: Total time = sum (e.g., 128ms for 20+5+3+100).
- Async: Total time ≈ longest task (100ms), more interrupts.

Device-Status Table:

- Managed by kernel.
- Each entry: contains type, address, state
- State updated after an interrupt (not function, idle, busy)

[3] Timer Interrupt

1. Interrupt occur after specific time period
2. Kills process / take back resource (e.g., unlock file)
3. In Unix: **clock routine**. Trigger by hardware per $\frac{1}{100}$ sec

System Calls

Goal: Abstract away hardware details. Provide interface for user.

Procedure

1. User process executes syscall
2. Look up **System Call Table**
3. CPU switch to kernel mode and execute syscall
4. Return status & results.

Types

Category	Functionalities	Windows	
Process control	<ul style="list-style-type: none"> - Create process, terminate process - Load, execute - Get process attributes, set process attributes - Wait event, signal event - Allocate and free memory 	CreateProcess(), ExitProcess(), WaitForSingleObject()	fork(), exit(), wait()
File management	<ul style="list-style-type: none"> - Create file, delete file - Open, close - Read, write, reposition - Get file attributes, set file attributes 	CreateFile(), ReadFile(), WriteFile(), CloseHandle()	open(), read(), write(), close()
Device management	<ul style="list-style-type: none"> - Request device, release device - Read, write, reposition - Get device attributes, set device attributes - Logically attach or detach devices 	SetConsoleMode(), ReadConsole(), WriteConsole()	ioctl(), read(), write()
Information maintenance	<ul style="list-style-type: none"> - Get time or date, set time or date - Get system data, set system data - Get process, file, or device attributes - Set process, file, or device attributes 	GetCurrentProcessID(), SetTimer(), Sleep()	getpid(), alarm(), sleep()
Communications	<ul style="list-style-type: none"> - Create, delete communication connection - Send, receive messages - Transfer status information - Attach or detach remote devices 	CreatePipe(), CreateFileMapping(), MapViewOfFile()	pipe(), shm_open(), mmap()
Protection	<ul style="list-style-type: none"> - Get file permissions - Set file permissions 	SetFileSecurity(), InitializeSecurityDescriptor(), SetSecurityDescriptorGroup()	chmod(), umask(), chown()

System Programs

1. File Management

- Functions: Create, delete, copy, rename, print, list, and manipulate files/directories.
- Example: `ls` command in UNIX lists directory contents.

2. Status Information

- Functions: Retrieve system data (e.g., date, time, memory, disk space, user count).
- Complex versions: Performance monitoring, logging, debugging.
- Output: Terminal, GUI, or files.
- Example: `top` command in Linux displays system resource usage.

3. File Modification

- Tools: text editors, searching command, text transformation files.
- Example: `vim` for editing files, `grep` for searching file contents.

4. Programming-Language Support

- Tools: Compilers, assemblers, debuggers, interpreters.
- Example: GCC for C/C++, Python interpreter.

5. Program Loading and Execution

- Tools: Loaders (absolute, relocatable), linkage editors, overlay loaders, debugging systems.
- Example: `gdb` for debugging C programs.

6. Communications

- Functions: Facilitate virtual connections between processes, users, and systems.
- Examples: Email, remote login (ssh), file transfer (scp).

7. Background Services (Daemons)

- Functions: Long-running processes for essential tasks (e.g., network connections, process scheduling, error monitoring).
- Examples: Network daemons, print servers.
- Note: Daemons often run at boot time and continue until the system shuts down.

Parameter Passing




Figure 2.7 Passing of parameters as a table.


1. **Registers:** fast, but limit space.
2. **Memory Block/Table:**
 - Put param in mem block
 - Put address of param mem blocks in registers
3. **Stack:** for variable-len data.

OS Types & Structure

Structure Type	Description	Examples	Pros	Cons
Monolithic	Single address space containing all OS functionality	Traditional UNIX, Linux base	High performance, efficient syscalls	Complex maintenance, tightly coupled
Layered	Hierarchical layers with defined interfaces; Layer N uses only layers 0 to N-1	-	Simple debugging, clear separation	Performance overhead, layer definition challenges
Microkernel	Minimal kernel (process/memory management + IPC); Services run as user processes	Mach (Darwin kernel component)	Extensible, reliable, secure	IPC overhead, complex coordination
Modular	Core kernel + Loadable Kernel Modules (LKMs); Dynamic loading/unloading of services	Linux modules, kernel extensions	Flexible, efficient communication	Potential reliability issues

Real-time: Processes have completion deadlines must be met

Distributed: Operate upon collection of computer across network



Type	Definition	Example of Use
Batch OS	Data or programs are collected grouped and processed at a later date.	Payroll, stock control and billing systems.
Interactive OS	Allows the user and the computer to be in direct two-way communication	Select from a menu at ATM.
Real-time OS	Inputs immediately affect the outputs. Timing is critical	Control of nuclear power plants, air traffic control systems.
Network OS	Allow a computer on a network to serve requests from other computers for data and provide access to other resources such as printer and file systems.	Manage simultaneous access by multiple users
Multiuuser OS	Handle many people running their programmes on the computer at the same time	A number of terminals communicating with a central computer which allocates processing time to each terminal in turn.
Multiprogramming OS	Ability to run many programmes apparently at the same time.	Mainframe systems. Each job is allocated a small amount of processing time (time slice) in the facilities.
Multitasking OS	The ability to hold several programmes in RAM at one time but the user switches between them.	Usually uses GUI's. Facilitates import and export of data.

I Hate Shell Script

ls usage

- `ls -l`: list in long format (detailed, verbose)
- `ls -a`: list all (include hidden files)
- `ls -R`: list subdirectories recursively
- `ls [hello]:` match single-character file h, e, l, o and o
- `ls {hello}*`: c: files starting with h, e, l, o and ending with .c
- `ls hello[1-3]*`: file starting with hello followed by 1, 2, or 3

Pipe: (`ls`; `cal`) | `wc -l` (`ls`; `cal`) creates a subshell; within this subshell, `ls` and `cal` run in two subprocesses (so 3 processes in total on the left hand side). Then, output piped to `wc` which runs in another

Unfamiliar Commands:

- `echo something\`: allow continue typing command on new line
- `echo $$`, `$$` evaluates to PID of current shell
- `wc`: outputs line count, word count (by space separate), **bytes** count, file name (if called with file name as parameter, eg. `wc test.txt`)
- `source`: execute a script or batch file
- `history`: list recent cmds; `!n` run command n; `!!` run last command
- `|` more: pipe output to more to display one screen at a time
- `export`: make variable available to subshells, but not parent shell
- ``cmd`` or `$(cmd)`: command substitution. ``` has undefined behavior
- `ps -ef` | `more` will display detailed information about all processes

Special Parameters

- `$*`: All positional parameters as a single string
- `$@`: All positional parameters as separate strings
- `$1`, `$2`, ...: Positional params (arguments to script), start from 1
- `$#`: Number of positional parameters (arguments)
- `$?`: Exit status of the most recently executed command
- `$$`: Process ID of the shell
- `$0`: Name of the shell or shell script
- `#!$#`: evaluates to the last argument

Exist Status: 0 for success, non-zero for failure, -1=255, 256=0

Access Modes: [dir or file] [user] [group] [other] -rwx r-- r--

- `chmod g+r f`: give group read access to file
- `chmod o-r f`: remove other's read access to file
- `chmod u+x f`: give user execute access to file
- `chmod -R g+r+w folder`: recursively give group read/write access to folder and all its content within.
- `chmod 777 f`: give all access to file
- Each digit represents user, group, other. E.g., 6 = (110)₂ = rw-
- **First character (before rwx, eg drwx-rwx-rwx):**
 - p: FIFO (named pipe)
 - -: Regular file
 - d: Directory
 - l: Symbolic link
 - s: Socket
 - c: Character device
 - b: Block device

\$PATH separated by : || To add current dir: PATH="\$PATH: ."

Regex:

- `^` means beginning of line;
- `$` means end of line;
- `.` means any character;
- `[abc]` any one of a, b, c;
- `[^abc]` anything except a, b, c;
- `\` use \ for escape.

grep Commands: `grep poly *.c` list all C files containing poly;
`grep -i poly *.c` same above, but case insensitive;
`grep -l -i 'poly' *.c` print only names of file with lines beginning with poly
`grep -h -i 'poly' *.c` print only lines beginning with poly

- Interpreted: PHP, Ruby, Perl, Bash, Lua, Tcl, Basic, HTML, JS, Py
- Compiled: COBOL, Fortran, Pascal, Ada, Lisp

Scripting Language:

- Advantages: fast to write, small size, partial execution, directly across OS services and execute commands, directly process output from other programs;
- Disadvantage: slow execution, weak data structure / typing, odd syntax error-prone, occasionally buggy

Random Concepts

- Run Commands together: `cmd1`; `cmd2`; `cmd3`
- Run cmd if previous success: `cmd1 && cmd2`
- Global variables by convention UPPERCASE, local lowercase
- For local variables use `local var=value`
- **DO NOT** add space before and after =
- **DO NOT** create a script called `test`, interfere with built-in `test` cmd
- `let` keyword
- `let x=a+b` is equivalent to `x=$((a+b))`
- but `x=a+b`, x is the literal string a+b
- **changes to the list** inside the loop will **not** affect the loop count

Special Option Variable

- **history**: enable command history, default = on
- **noclobber**: prevent overwrite file when I/O redirection, default = off

- **First Come First Serve (FCFS):** First process in, first process out.
- **Shortest Job First (SJF):** Shortest job first.
- **Multi Level Queue:** Ready queue divided into multiple queues.
 - Each queue has own scheduling algorithm
 - Scheduling **between queues**.
 - Fixed Priority Scheduling: each queue has own priority
 - Time Slicing: each queue gets a time slice (like RR)



Preemptive:

- **Shortest Remaining Time (SRT):** Preemptive version of SJF.
- **Round Robin (RR):**
 - Predefined *time quantum* (e.g., $q = 10\text{ms}$).
- **Priority Scheduling:**
 - *Priority Number* assigned to each process. Higher priority first.
 - SJF is priority scheduling where priority is the time to completion.
 - FCFS is priority scheduling where priority is the time of arrival.

Strengths & Weaknesses

- SJF / SRT: Optimal for minimizing waiting time, good for batch jobs.
 - SJF proven to have lower TAT, **optimal** for non-preemptive
 - SRT proven to have lower TAT, **optimal** for preemptive
- RR: fair to all process, good for interactive jobs.
- FCFS: simple, no starvation since all will eventually execute. But *Convoy Effect*, where short process stuck behind long process.

Common Scheduling Issues & Solutions

Issue 1: Starvation

- **Problem:** Low-priority processes may never execute due to continuous arrival of high-priority processes
- **Solution:** Implement *aging* mechanism - processes gradually increase in priority the longer they wait

Issue 2: High Overhead

- **Problem:** Complex scheduling algorithms (e.g., multilevel feedback queues) create significant system overhead
- **Solution:** Optimize queue management with efficient data structures and algorithms

Issue 3: Priority Inversion

- **Problem:** Higher priority process waits while lower priority process holds a needed resource
- **Solution:** Implement *priority inheritance* - lower priority process temporarily inherits the higher priority

Issue 4: Unpredictability

- **Problem:** Execution order can be unpredictable, problematic for real-time systems requiring timing guarantees
- **Solution:** Use deterministic scheduling algorithms or provide explicit timing guarantees for real-time requirements

Issue 5: Parameter Tuning

- **Problem:** Selecting optimal values (time quantum, priority levels) significantly impacts performance
- **Solution:** Systematic testing and performance analysis to determine optimal parameters

Memory Management

- **Logical Address:** virtual address, generated by CPU
- **Physical Address:** absolute address, address sent to memory units.

Binding: logical \rightarrow physical addr translation

- **Compile-time:** Phys addr hardcoded. Recompile if program moved.
- **Load-time:** *Relocatable code* (relative addr). Phys addr calculated when loaded to RAM; address fixed afterwards.
- **Execution-time (Run-time):** Bind delayed till run. Process can move during execution. Logical addr used; MMU calc phys addr dynamically. Needs HW support (MMU).

Memory Management Unit (MMU)

- **Relocation Register (RR):** added to every logical addr to get base.
- **Limit Register (LR):** offset limit from the base physical memory

Contiguous Mem. Allocation:

- Entire process in single block of mem
- Multiple fixed partition: multiprogramming w/ fixed num. of tasks
- Mult. variable partition: multiprogramming w/ varying num. of tasks
 - Process arrives \rightarrow finds **hole** large enough to fit it
 - First-fit: find first hole that's big enough.
 - Best-fit: find smallest hole big enough (little leftovers, not useful)
 - Worst-fit: find largest hole (large but useful leftovers)
- **First-fit and Best-fit** are normally better
- **External Fragmentation:** space exist, but not contiguous (fixed with *compaction* - stop all process and group memory together)
- **Internal Fragmentation:** waste within partition (hole slightly larger than process, and managing this overhead is not worth it)

Non-Cont. Mem. Alloc:

- multiple blocks (*paging* or *segmentation*)
- **Paging:** physical mem divided into fixed-size **frames**. Logical mem divided into same fixed-size **pages**. Page table maps page \rightarrow frames
 - logical addr. = page number + offset
 - *Page number:* index of page table. *Offset:* offset within page.
 - **Paging suffers Internal Fragmentation**
 - **Paging allows sharing a program across multiple processes**
- **Implementation:**
 - *page-table base register* (PTBR): points to page table of a process in the memory. **Fast context switch:** just change PTBR.
 - But needs **two mem. accesses**: 1. get page table, 2. get the data.
 - **Solution:** *Translation Look-aside Buffer* (TLB) - store recently translated logical page numbers.
 - cache access time c , memory access time m , TLB hit ratio h
 - Translation time = $h \times c + (1 - h) \times (c + m)$
 - Effective access time = $c + (2 - h) \times m$
- **Hierarchical Paging** (otherwise, page table too large):
 - Example, 32bit machine - 20bit for page number, 12bit for offset.
 - Page the page table: 10 bit page number, 10 bit offset
- **Memory Protection:**
 - *Valid-Invalid bit:* additional bit to each page table entry
 - *Page-table length register* (PTLR): check validity of logical addr.

- **Segmentation:** logical memory space is different-sized segments. Each segment table entry contains *base*, *limit*, *valid bit*
 - Segment-table base register (STBR): point to segment table base in memory. *Segment-table length register* (STLR): number of segment
 - **Segmentation suffers External Fragmentation**
 - **Segmentation also allows sharing. Put same entry in each table**

Virtual Memory Techniques

- **Overlay:** break program into stages, and load them sequentially
 - swapping but on different **stages** instead of parts of program
 - Memory has a common area for common modules
 - When program $>$ physical memory (eg embedded systems)
- **Vritual memory:** store process on disk, load on demand.
 - Uses **memory map** (basically page table but for virtual mem)
 - More general puposed

Checking memory usage

- `a.out & runs a.out` in background. **allow multiple instance**
- Then `ps -l` lists current running process in this terminal.
- Columns: VSZ total size. RSS mem usage ($VSZ \geq RSS$)

Methods:

- **Demand paging:** load page only when it's needed.
 - Need for a page indicated by PC or MAR.
- Via lazy process **swapper**. Swapper that swaps pages = **pager**.
- Valid-invalid bit: indicate whether page is in memory
 - **Procedure:** 1. check valid-invalid bit. 2. if invalid, trigger a *trap* called *page fault*. 3. Run *page fault handler*, that ask *pager* to swap
 - Note: valid-invalid bit set to invalid for all entries initially.
 - **Page fault handling:** 1. get empty frame from free frame list. 2. load page from disk. 3. update page table to point to new frame. 4. set valid-invalid bit to valid. 5. Restart page fault instruction.
- **Performance:**
 - *page fault rate* = $0 \leq f \leq 1$, page fault service time = page fault overhead + time to swap
 - *memory access time* = m , fault overhead + time to swap
 - *page fault service time* = s , page in + restart overhead
 - *Effective Access Time (EAT): expected memory access time* $EAT = (1 - f) \times m + f \times s$
- **Anticipatory paging:** predict what page will be needed next.

Page Replacement

- **Belady's Anomaly:** page fault rate \uparrow if number of page frames \uparrow
- **Reference String:** page number sequence
- **FIFO:** Always replace oldest page. Suffer from *Belady's Anomaly*.
- **Optimal:** Replace page that will be used furthest in the future.
- **LRU:** Replace page that has not been used for the longest time.
- **Allocation:**
 - *Local page replacement:* process only replaces its own frames
 - *Global page replacement:* process can replace other's frames
 - **Fixed allocation (only local page replacement allowed):**
 - Equal allocation: each process gets same number of frames ($\frac{P}{n}$)
 - Proportional allocation: each get frames proportion to its size
 - *Variable allocation:* number of frames vary over time. Give processes with too many page faults more frames.

Thrashing: when number of available frames $<$ total size of active pages of the processes in the ready queue.


- Solution: **reduce** the degree of multiprogramming.
- **Working-Set:** total number of unique pages reference in a period Δ
- When working-set size $>$ number of frames, thrashing occurs
- **Page Fault Frequency (PFF) algorithm:** define acceptable range. If rate too low, take away some frames. If rate too high, add frames.

Deadlock

- **Deadlock:** set of blocked resources. cannot proceed. **Conditions:**
 - *Mutual Exclusion:* only one process can use a resource at a time
 - *Hold and wait:* process holds ≥ 1 resource, waits another resource
 - *No preemption:* resource only released voluntarily by process
 - *Circular wait:* hold-and-wait processes forms a cvcle.
- **Livelock:** can be resolved by chance.

Resource Allocation Graph (RAG)

- $P = \{P_1, P_2, \dots, P_n\}$, set of all processes
- $R = \{R_1, R_2, \dots, R_m\}$, set of all resources
- **Request Edge:** P_i requests R_j ($P_i \rightarrow R_j$)
- **Assignment Edge:** $R_j \rightarrow P_i$
- If a req can fulfill, **request edge becomes assignment edge**



Check for deadlock (RULES):

- If no cycle, **definitely no deadlock**
- If cycle, only one instance per resource type, **deadlock**
- If cycle, multiple instances per resource type, **might deadlock**

Deadlock Handling

- **Ostrich approach (Deadlock ignorance):** ignore the deadlock
- **Deadlock avoidance (Banker's algorithm):** don't enter deadlock state
- **Deadlock prevention:** ensure never enter deadlock state
- **Deadlock detection:** detect deadlock after happen; recover from it

Deadlock Prevention:

- Ensure when process request resource, it is not holding other resources (**in-degree/out-degree cannot be non-zero at the same time**)
- Make preemption possible (e.g., by using time quantum)
- Prevent circular wait:
 - Order all resource types with a pre-defined number
 - P_i can only request R_j if j larger than the number of any resource currently held by P_i . Else P_i have to release larger resource first
 - (**All edge point from smaller number R to large. No reverse edges**)

Deadlock Avoidance (Banker's algorithm):

- **Safe State:** system is safe if exist certain resource allocation order that can avoid deadlock. This order is called **Safe Sequence**
- $\{P_1, P_2, \dots, P_n\}$ is safe sequence if: for each P_i , its *need* can be satisfied by current available resource + resources held by all $P_{j < i}$

Deadlock Detection (Wait-for Graph)

- Detect deadlock by *deadlock detection algorithm*
- Execute *deadlock recovery scheme* to recover from deadlock
 - *Process termination:* kill all deadlocked processes / kill one victim process at a time until deadlock eliminated (consider: priority progress, etc.) Might **starvation**: some process always victim
 - *Process rollback:* return to safe state and retry

Banker's Algorithm

- Avail $[m]$ # of resources m currently available
- Alloc $_i[m]$ # of resources m currently allocated to process i
- Max $_i[m]$ # of resources m that process i will at most need
- Need $_i[m]$ Max $_i[m]$ - Alloc $_i[m]$ (how much more i needs)

Intuition:

1. Simulate allocate resource to a P_k
 2. Assume P_k run in background (async)
 3. **If** (synchronized) execution sequence exists then **safe** Else **unsafe**
- | | |
|--|---|
| 1 Work \leftarrow Avail (array of length m) | 2 Finish \leftarrow [False, ...] (array of length n) |
| 3 while not done do: | 4 find i s.t. \neg Finish $[i]$ and Need $_i \leq$ Work |
| 5 if no such i exists, break | 6 Work \leftarrow Work + Alloc $_i$ |
| 7 Finish $[i] \leftarrow$ True | 8 if Finish $[i] =$ False $\forall i$ then return <i>Safe</i> |
| 9 else return <i>Unsafe</i> | |

Find all safe sequence: requires drawing search (DAG) graph.

Wait-for Graph

RAG except simplify **FROM** $P_i \rightarrow R_j$ and $R_j \rightarrow P_i$ **TO** $P_i \rightarrow P_j$ Periodically check for cycles in the *Wait-for Graph*

Synchronization

Critical Section (CS) Problem:

- **Critical Section:** code segment where shared data is accessed
- Solution: [Entry Section] + [Critical Section] + [Exit Section]
- **Rules:**
 1. *Mutual Exclusion:* only one process in their CS at a time
 2. *Progress:* When no one in CS, and one process wants to enter CS decision must be made within finite time. (**no hanging request!**)
 3. *Bounded Waiting:* Between request and grant for A, limit the number of times B enters CS. (**fairness / no starvation!**)
- **Solutions:**

- **Lock Variable:** shared boolean

```
while (lock == true) wait();
lock = true; Enter_CS(); lock = false;
```

Problem:

1. Process P_0 checks lock; it's false.
2. Context switch **before** P_0 can set lock = true.
3. Process P_1 checks lock; it's still false.
4. P_1 sets lock = true and enters its critical section.
5. Context switch back to P_0 (already passed while). Sets lock = true (even though it's already true) and also enters its CS.
6. Now both P_0 and P_1 are in their CS simultaneously!!!

- **Test-and-Set (TSL) Lock:** use special hardware instruction, TSL. This instruction is atomic. **It is a single, indivisible operation**

```
while(TestAndSet(&lock) == true) wait();
Enter_CS(); lock = false;
```

- **Turn Variable:** two processes P_0 and P_1 . A shared integer variable **turn** to indicate **whose turn it is to enter the critical section**

```
while(turn != self) wait();
Enter_CS(); turn = other;
```

Solution	Mutual Excl.	Progress	Bounded Wait
Lock Variable	Fails	Fails	Fails
Test and Set	Satisfied	Satisfied	Fails
Turn Variable	Satisfied	Fails	Satisfied

Peterson's Algorithm

Peterson's Algorithm (flag helps satisfy *Progress*)

- 1 turn $\leftarrow 0$ // indicate whose turn to enter CS
- 2 flag[0] \leftarrow flag[1] \leftarrow false // whether process ready to enter CS
- 3 **while** true **do:**
 - 4 flag[i] \leftarrow true
 - 5 turn $\leftarrow j$ // Use j to indicate other processes ($j \neq i$)
 - 6 **while** flag[j] = true and turn = j **do** wait() **end**
 - 7 **[Critical Section]**
 - 8 flag[i] \leftarrow false
 - 9 **Remainder Section ...**

Pros: Builds on top of **Turn Variable** and satisfies *Progress*

Cons: while-loop waste CPU time. The thing it's waiting starved.

Semaphores

- Ultimate solution for critical section problem

```
while (true) {
    P(S);
    Enter_CS();
    V(S);
    Remainder_Section();
}
S is a semaphore, shared by all processes
P(S): decrease S by 1, if negative move current process to waiting queue
V(S): increase S by 1, if exist processes in waiting queue, wake up one of them
```

Intuition:

- **S** is a counter of available resources
- **P(S)** check if any resource available. If true, take one ($S -= 1$). Otherwise ($S \leq 0$), this process sent to waiting queue
- **V(S):** called when process release a resource ($S += 1$). If there is any process in waiting queue, wake up one of them.
- Both **P(S)** and **V(S)** are **atomic operations**. Can't context switch.

P(S), aka: wait(), down(), acquire()
- V(S), aka: signal(), up(), release()

Other Concepts

- Counting Semaphore:** $[-\infty, +\infty]$, though typically non-negative.
 - For resource that has multiple instances
- Binary Semaphore (Mutex):** $[0, 1]$, Init. to 1, alter between 0 and 1.
 - For resource that has only one instance.

Producer-Consumer Solution (using Semaphores)

```
// producer // consumer
while (true) while (true)
{ {
  down(empty); down(full);
  down(mutex); down(mutex);
  produce(); consume();
  up(mutex); up(mutex);
  up(full); up(empty);
} }
```

empty: # of empty buffer slots
full: # of occupied slots
mutex: ensure only one party access the buffer at a time

If wait(mutex) placed before wait(empty), producer can lock buffer and do nothing (if buffer is full) cause **deadlock**

File System, Secondary Storage

Access Methods

- Sequential Access:** Data accessed in order, beginning to end
 - read next:** read next data, advance pointer (read fp; fp++)
 - write next:** write next data, advance pointer (write fp; fp++)
 - reset:** put file pointer back to the beginning (fp=0)
 - skip forward/back:** fp forward/backward without reading.
- Direct Access:** aka Relative access. Files are fixed-length *records*. Direct access like arrays, but based on *record number*
 - Read n:** return n-th data item / block
 - Write n:** update n-th data item / block
- Indexed Access:** Use a separate *index file* (aka *direct file* or *index file*) contain pointers to data blocks of data file. Can direct access

Allocation Methods

Methods should 1) Efficient disk util. 2) fast access time

- Contiguous Allocation:** Each file stored as a contiguous block
 - Pro:** easy to implement, fast access time (sequential / direct)
 - Con:** Disk fragmented, difficult to grow file
- Linked List Allocation:** File has many blocks, linked by pointer
 - Pro:** No external fragmentation. File size can increase
 - Con:** Large seek time, direct access hard, pointer overhead
- Indexed Allocation:** index file points to each file's first block
 - Pro:** No external fragmentation. Supports direct access
 - Cons:** Pointer (memory) overhead. Multi-level index (if file big)

UNIX I-Node Structure

Attributes
Direct Blocks
Single Indirect Blocks
Double Indirect Blocks
Triple Indirect Blocks

Disk Scheduling Algorithm

- Goal: Minimize seek time
- Seek time = time taken to reach desired track

FCFS (First Come First Serve)
SSTF (Shortest Seek Time First)

- Greedily choose the closest request from current head position.
- Issue: starvation problem. Overhead calculating seek time.

C-SCAN (Circular SCAN)

- Move from current to the largest track (read along the way)
- Then immediately back to smallest **requested** track (**don't read**)
- Read from smallest track to last requested track (read along way)

SCAN (assume larger value indicator, otherwise vice versa)

- Minimize seek time using *Elevator algorithm*
- 1. Move head to the largest track (while reading along the way)
- 2. Then back to the smallest **requested** track (read along way)

SCAN better than C-SCAN

LOOK (assume direction towards large value, if not then vice versa)

- Move from current to largest **requested** track (read along way)
- Move back to smallest **requested** track (read along the way)

- LOOK better than SCAN**

C-LOOK (Circular LOOK)

- Move from current to largest **requested** track (read along way)
- Immediately back to smallest **requested** track (don't read)
- Then read to final requested track (read along way)

Redundant Arrays of Inexpensive Disks (RAID)

Level	Description	Util. (N=#disk)
RAID0	Stripe. No Fault tolerance	100%(N/N)
RAID1	Mirror (store 2 copies)	50%(N/2)
RAID2	Bit-lvl stripe + ECC	Low, (N - P)N
RAID3	Byte-lvl stripe. dedicate parity disk	(N - 1)/N
RAID4	Block-lvl stripe. dedicate parity disk	(N - 1)/N
RAID5	Block-lvl. distributed parity disk	(N - 1)/N

P (ECC disks) must satisfy: $2^P \geq N + P + 1$

Protection

Principle of Least Privilege:
Programs, user, systems given **just enough** to perform their task
Limit damage if bug or get abused

Separate *policy* from *mechanism*:

- Mechanism:** the actual implementation that enforces the *policy*
- Policy:** customizable security rules realized through *mechanism*

Domain specifies a set of object and their operations

- Access-right = <object-name, right-set>

Access Matrix

- Row represent *domains*
- Column represent *objects*
- Access(i, j) tells the set of operation Domain_i can perform on Object_j

	object	F ₁	F ₂	F ₃	printer
domain	D ₁	read		read	
D ₂					print
D ₃			read		execute
D ₄		read	write	read	write

- Entire table: stored as **ordered triples**
 - Result: List of triples <domain, object, right>
- Each column called **access control list**
 - Result: per-object list of ordered pairs: <domain, right>
- Each row called **capability list**

```
fill() {
  local size=$1 content=$2 blob=()
  for ((i=0; i<size; i++)); do blob+=("$content"); done
  echo "${blob[@]}"
}

lookup() {
  local mode=$1 query=$2 index_only=$3 keys vals i results=()
  if [[ $mode == "name" ]]; then
    keys=("${emp_nms[@]}")
    vals=("${emp_ids[@]}")
  else
    keys=("${emp_ids[@]}")
    vals=("${emp_nms[@]}")
  fi
  for ((i=0; i<total_emps; i++)); do
    if [[ ${keys[i]} == $query ]]; then
      if (( index_only )); then results+=("$i");
      else results+=("${vals[i]}"); fi
    fi
  done
  echo "${results[@]}"
}

push_unique() {
  local i
  for i in "${queue[@]"; do
    ((i == $1)) && return
  done
  queue+=($1)
}
```

```
check() {
  [[ -e "$1" ]] && [[ -f "$1" ]] && [[ -r "$1" ]] || { echo "Invalid
File: $1"; exit 1; }
}

parse() {
  local flag=0
  for p in "${params[@]"; do
    [[ $p == "employee" ]] && { flag=1; continue; }
    if (( flag )); then
      emps+=("$p")
    else
      [[ " ${dats[*]} " != *" $p " ]] && dats+=("$p")
    fi
  done
  n_emps=${#emps[@]}
}

rd() {
  # read employee
  while IFS= read -r line || [[ -n "$line" ]]; do
    read id nm <<< "$line"; emp_ids+=("$id"); emp_nms+=("$nm")
    done < "employees.dat"
    total_emps=${#emp_ids[@]} # total number of employees

  # read data files
  for file in "${dats[@]"; do
    check "$file" # check file validity
    local firstline=1 index=1 date=0
    while IFS= read -r line || [[ -n "$line" ]]; do
      if (( firstline )); then
        firstline=0
        local tmp1 dpt month year; read tmp1 dpt month year <<< "$line"
        date=$(parse_date "$year" "$month")
        # find the index of the department
        local i
        for (( i=0; i<${#departments[@]}; i++ )); do
          if [[ ${departments[i]} == $dpt ]]; then
            index=$i; break; fi
          done
          # if department not found, add the department name
          # then expand the attendance status array size
          # then initialize new slots with "N/A"
          if (( index == -1 )); then
            departments+=("$dpt")
            status+=($(fill "$total_emps" "N/A"))
            last_modified+=($(fill "$total_emps" 0))
            index=$(( ${#departments[@]} - 1 ))
          fi
          else
            local id attd; read id attd <<< "$line"
            local i=${lookup "id" "$id" 1} # look up the index
            # of ID
            local j=$(( index * total_emps + i ))
            if (( date >= ${last_modified[j]} )); then
              last_modified[j]=$date; status[j]=$attd; fi
              fi
              done < "$file"
            done

  # echo results
  out() {
    # remove redundancy
    for e in "${emps[@]"; do
      # combine indices from both name and id searches in
      one array
      local found_indices=("${lookup "name" "$e" 1} ${lookup "id"
"$e" 1})
      if (( ${#found_indices[@]} == 0 )); then
        echo "Invalid Entry: $e"
        exit 1
      fi
      for index in "${found_indices[@]"; do
        push_unique "$index"
      done
    done

  # loop through employees' indices
  for ei in "${queue[@]"; do
    echo "Attendance Report for ${emp_ids[ei]} ${emp_nms[ei]}"
    local total_present=0 total_absent=0 total_leave=0

  # loop through known departments
  for di in "${!departments[@]"; do
    local j=$(( di * total_emps + ei ))
    local department_name=${departments[di]}
    local department_status=${status[j]}

    if [[ $department_status == "N/A" ]]; then
      echo "Department ${department_name}:
$department_status"
    else
      echo "Department ${department_name}:
$department_status 1"
      fi
      # increment counters based on status
      [[ $department_status == "Present" ]] && ((total_present++))
      [[ $department_status == "Absent" ]] && ((total_absent++))
      [[ $department_status == "Leave" ]] && ((total_leave++))
    done
    echo

  # output statistics
  echo "Total Present Days: $total_present"
  echo "Total Absent Days: $total_absent"
  echo "Total Leave Days: $total_leave"
  done
}

main() {
  parse
  rd
  out
}
```