

Measurement Error

- Error Rate: # of errors / # of instances = $(FN+FP)/N$
- Recall: # found positives / # actual positives = $TP/(TP+FN)$
- Precision: # found positives / # predicted positives = $TP/(TP+FP)$
- Specificity = $TN/(TN+FP)$
- False Alarm Rate = $FP/(FP+TN) = 1 - \text{Specificity}$

Association Rule Mining

For finding frequent patterns, associations, corr., causal structures.

Format: $X \rightarrow Y$ [support, confidence]

- Support:** $P(X \cap Y)$
 - prob. that transaction contains both X & Y
 - indicates statistical independence of association rule
- Confidence:** $P(Y|X) = \frac{P(X \cap Y)}{P(X)} = \frac{P(X|Y)P(Y)}{P(X)}$
 - prob. that transaction having X also has Y
 - the degree of correlation. Measures rule's strength

Apriori Algorithm

```
# C[k]: Candidate itemset of size k
# L[k]: Frequent itemset of size k
L[1] = [frequent single items]; k = 2
while L[k-1] is not empty:
    C[k] = generate_candidates_from(L[k-1])
    C[k] = prune_candidates(C[k])
    L[k] = count_support(C[k])
```

Step 1: Generate Candidates

Generate candidates for k -itemset from $(k-1)$ -itemset:

```
n = len(L[k-1])
for i in range(n):
    for j in range(i+1, n):
        itemset_i, itemset_j = L[k-1][i], L[k-1][j]
        shared_items = intersection(itemset_i, itemset_j)
        if len(shared_items) == k-2:
            candidate = join(itemset_i, itemset_j)
            C[k].append(candidate)
```

Step 2: Prune Candidates

Prune candidates that do not meet the minimum support threshold:

```
# given k-itemset, prune if any (k-1) subset not in L[k-1]
for cand in C[k]:
    for split in range(len(cand)):
        subset = cand[:split] + cand[split+1:]
        if subset not in L[k-1]:
            C[k].remove(cand)
            break
```

Step 3: Counting

```
cnt = {cand: 0 for cand in C[k]}
for entry in database:
    for cand in C[k]:
        if cand in entry:
            cnt[cand] += 1
L[k] = [cand for cand in C[k] if cnt[cand] >= min_support]
```

Repeat Step 1 & 2 until $L[k]$ is empty

Rule Generation

- For each frequent itemset L , gen all non-empty subsets of L
- For each non-empty subset $s \in L$, gen rule $s \Rightarrow (L - s)$

Other Measures

Objective Measures: support & confidence

Subjective Measures: Interest (lift)

$$\text{Lift} = \frac{\text{Conf}(X \Rightarrow Y)}{P(Y)} = \frac{P(X \cap Y)}{P(X)P(Y)}$$

Intuition 1: Performace (above baseline) Test

- $\text{Conf}(X \Rightarrow Y)$ given X happened, prob. of Y also happening; $P(Y)$, prob. of Y happening in general (**baseline**)

Intuition 2: Independence Test

- $P(X \cap Y)$ **actual observed** prob. of $X \& Y$ happening together;
- $P(X)P(Y)$ expected prob. of $X \& Y$ happening together if they were completely independent. (**the independence test**)

Sequential Rule Mining

Step1-Sort Phase: First group data by customer ID. For each customer, sort their transaction by timestamp.

Step2-Freq Itemset: Find all freq/large itemset L (using Apriori).

- Counting rule: number of customers whose sequence contains the itemset at least once

Step3-Transform: Transform each sequence to frequent itemset

- Given: Freq. Itemsets (from Step 2): $\{(A), (B), (A, B), (C)\}$
- Trans: $< (A, B), (C), (D) > \rightarrow < \{(A), (B), (A, B), \{(C)\} >$
- (A, B) became $\{(A), (B), (A, B)\}$
- (C) became $\{(C)\}$
- (D) is removed, because it is not in Freq. Itemsets

Step4-Sequence: Find sequences using frequent itemsets (AprioriAll)

- Given: transformed seq (from step 3) and freq 1-itemset (from step 2)
- Find frequent k -sequences, by joining $(k-1)$ -sequences with the same $(k-2)$ prefix. Prune if any $(k-1)$ -subseq is not frequent.
- Count support for each k -sequence, and remove if below threshold.
- Given candidate seq ABC, transaction $< (A, B), (C), (D) >$ supports ABC; transformed version helps make this counting easier: $< \{(A), (B), (A, B)\}, \{(C)\} >$ **This is why Step3 is needed!!!**

Step5-Maximal Phase:

Sequence s is maximal if s is not contained in any other sequences

```
for k in range(max_k, 1, -1):
    for seq in L[k]: for subseq in seq:
        if subseq in L[k-1]: L[k-1].remove(subseq)
```

AprioriAll Alogrithm

```
L[1] = [frequent 1-itemsets]; k = 2
while L[k-1] is not empty:
    C[k] = generate_candidates_from(L[k-1])
    L[k] = prune_candidates(C[k])
```

Diff. compared to Apriori

generate_candidates_from:

- given two $(k-1)$ -sequences A and B , generate k -sequence C ,
- if $A[-1] == B[-1]$, we get: $A + [B[-1]]$ and $B + [A[-1]]$
- A and B must share first $(k-2)$ items!
- Example:** permute last item of $(k-1)$ -sequence
 - $< A, B > + < A, C > \Rightarrow < A, B, C >$ and $< A, C, B >$

prune_candidates:

- given a k -seq, prune if **any** of its $(k-1)$ -subseqs is not frequent.
- Example:** $L[3] = \{\text{abc, abd}\}$, generate abdc and abcd
 - Prune abdc because adc and bdc are not in $L[3]$

Decision Tree (DT)

Two Stages

- Tree construction
- Tree pruning (rm branch that reflects outliers and noise)

Tree Construction

ID3/C4.5: information gain

- Attributes assumed categorical
- Can be modified for continuous-valued attributes

IBM IntelligentMiner: gini index

- Attributes assumed continuous
- Need other tools to get possible split values
- Can be modified for categorical attributes

Information Gain

Definition:

- How much did information entropy decrease by?
- Information Entropy \equiv impurity
- $\text{Gain}(D, a)$: Information gain if we were to split samples D by attribute a into V subsets $\{D_1, D_2, \dots, D_V\}$.

Equations

Information : $I(x_i) = -\log_2(p(x_i))$

Information Entropy : $\text{Ent}(X) = E(I(X))$

$$\begin{aligned} &= \sum_{i=1}^n p(x_i)I(x_i) \\ &= \sum_{i=1}^n p(x_i)(-\log_2(p(x_i))) \\ &= -\sum_{i=1}^n p(x_i)(\log_2(p(x_i))) \end{aligned}$$

Information Gain : $\text{Gain}(D, a) = \text{Ent}(D) - \sum_{v=1}^V \frac{|D_v|}{|D|} \text{Ent}(D_v)$

Impurity After Split : $= -\sum_{v=1}^V \frac{|D_v|}{|D|} \sum_{x_i \in D_v} p(x_i) \log_2 p(x_i)$

Everything Together:

$$\begin{aligned} \text{Gain}(D, a) &= \\ &= \text{Ent}(D) - \sum_{v=1}^V \frac{|D_v|}{|D|} \text{Ent}(D_v) \\ &= -\left(\sum_{i=1}^n p(x_i) \log p(x_i) - \sum_{v=1}^V \frac{|D_v|}{|D|} \sum_{x_i \in D_v} p(x_i) \log p(x_i) \right) \end{aligned}$$

- For categorical values V is the number of distinct values of a
- $|D_v|$ is the size v 'th split. $|D|$ is the size of the entire dataset

Regression Tree

- Prediction is computed as average of sample values in the leaf node.
- Impurity measured as sum of squared deviations from leaf mean
- Performance measured by root mean squared error (RMSE)

Tree Pruning & Enhancements

Pre-pruning:

- Halt tree construction early
- Halt before goodness measure falls below certain threshold.
- (hard to choose appropriate threshold)

Post-pruning:

- Removes branches from a "fully grown" tree.
- Progressively remove branches.
- Use testing set to decide the best pruned tree.

Other Enhancements

Continuous-valued Attributes:

- Treat each unique continuous attribute value in the dataset as discrete and identify all possible split points between them.
- e.g. $[21.6, 22.0, 23.0]$. Try splits at: $[21.8, 22.5]$
- Time Complexity: $O(N)$ after sort. $O(N \log N)$ in total.


Missing Values:

- Method 1: Assign most common value of that attribute
- Method 2: Assign probability to each of the possible values

Attribute Construction:

- Create new attribute on existing ones (e.g., spare ones)
- Purpose: Reduce fragmentation, repetition, replication

Non-axis Parallel Split



Divide & Conquer

Internal Decision Nodes:

- Univariate: Use single attribute a_i
 - Numeric: Binary split
 - Discrete n -way split, for all n possible values
- Multivariate: use all attributes $a_i \in A$

Leaves:

- Classification: class labels
- Regression: Numeric - avg of leaf nodes / local fit (fit a more sophisticated model locally, e.g., linear regression)

Pseudocode

```
def GenerateTree(x):
    if NodeEntropy(x) < 0.1 /* (See I_m on slide 29)
        Create leaf labelled by majority class in x
        return
    i ← SplitAttribute(x)
    for each branch of x[i]:
        Find x[i] falling in branch
        GenerateTree(x[i])

def SplitAttribute(x):
    MinEnt ← MAX
    for all attributes i = 1, ..., d:
        if x[i] is discrete with n values:
            Split x into x[1], ..., x[n] on x[i]
            e ← SplitEntropy(x[1], ..., x[n])
            if e < MinEnt:
                MinEnt ← e
                bestf ← i
        else: # x[i] is numeric
            for all possible splits
                Split x into x[1], x[2] on x[i]
                e ← SplitEntropy(x[1], x[2])
                if e < MinEnt:
                    MinEnt ← e
                    bestf ← i
    return bestf
```

Exam Tips (Rule Extraction Format)

- R1: IF (age > 38.5) AND (years-in-job > 2.5) THEN y = 0.8
 R2: IF (age > 38.5) AND (years-in-job ≤ 2.5) THEN y = 0.6
 R3: IF (age ≤ 38.5) AND (job-type = 'A') THEN y = 0.4
 R4: IF (age ≤ 38.5) AND (job-type = 'B') THEN y = 0.3
 R5: IF (age ≤ 38.5) AND (job-type = 'C') THEN y = 0.2

Bayes Classification

Given sample X and class C :

$$P(C|X) = \frac{P(X|C)P(C)}{P(X)} \text{ aka. posterior } = \frac{\text{prior} \times \text{likelihood}}{\text{evidence}}$$

Idea: assign sample X the best C such that $P(C|X)$ is maximized.

- $P(X)$ unknown, but constant for all C . So we can ignore it.
- $P(C)$ unknown, but reasonably estimated from dataset
- $P(X|C)$ unknown and infeasible.
 - Need to know the joint probability $P(x_1, x_2, \dots, x_n | C)$.
 - If each x_i takes 2 values, we can have 2^n possible combinations.
 - Probability that any such sample already exist in our dataset ≈ 0

Naive Bayes Classifier

Assumes attributes are conditionally independent:

$$P(C|X) \propto P(C) \prod_{i=1}^n P(x_i | C)$$

How to estimate $P(x_i | C)$?

- If x_i categorical: relative frequency of x_i in class C in dataset
- If x_i continuous: gaussian density via all i -th attr. in class C

KNN (K-Nearest Neighbors)

Instance-Based Methods (a.k.a memory-based)

- Instance-based classification: Store training examples, delay processing until a new instance must be classified ("Lazy Evaluation")

Examples:

- k -nearest neighbors: instances as points in Euclidean space.
- Case-based reasoning: symbolic representation and knowledge-based inference

Procedure

- Time Complexity: $O(n \log k)$ via min-heap of size k

- Compute distance to all other training records.
- Sort the distance, and find k nearest neighbors
- Use the k nearest neighbors to determine the class label of the unknown sample (e.g., through majority voting)
- Increasing k : Less sensitive to noise
- Decreasing k : Capture finer structure

Distance Metrics

- Continuous Attributes: Euclidean Distance

- normalize each dimension by standard deviation

• Discrete Data: use hamming distance

Curse of Dimensionality

Definition: Prediction accuracy degrade quickly when number of attributes grow, because

- When many irrelevant attributes involved, relevant attributes are shadowed, and distance measurements are unreliable

Solution:

- Remove irrelevant attributes in pre-processing (e.g., dimensionality reduction such as PCA)

Ensemble Methods

Advantages: Improved accuracy; other types of classifiers can be directly included; Easy to implement; Not too much parameter tuning

Disadvantage: Black box; Not compact representation.

Who Minimize Variance:

- Bagging
- Random Forest

Who Minimize Bias:

- Boosting
- Ensemble Selection

Proof

Consider binary classification problem $y \in \{-1, +1\}$, with ground truth function f , base classifier h_i , and base classifier error rate ε :

- Base Classifier Error Rate: $P(h_{i(x)} \neq f(x)) = \varepsilon$
- Ensemble Classifier: $H(x) = \text{sign}(\sum_i^T h_{i(x)})$

Assume independence of h_i error rates. From *Hoeffding Inequality*:

$$P(H(x) \neq f(x)) = \sum_{k=0}^{[T/2]} \binom{T}{k} (1-\varepsilon)^k \varepsilon^{T-k} \leq \exp\left(-\frac{1}{2}T(1-2\varepsilon)^2\right)$$

Therefore, as the number of base classifiers T increase, the probability of misclassification decrease exponentially

Bagging

Bagging = Bootstrap Aggregation

Traditional Bagging

- Takes original dataset D with N training examples
- Creates M different copies $\{D_m\}_{m=1}^M$ via
 - Sampling from D with replacement
 - Each D_m has same number of examples as D
- Train base classifiers h_1, \dots, h_M using D_1, \dots, D_m
- Average / Vote model as the final ensemble $h = \frac{1}{M} \sum_{m=1}^M h_m$

Random Forest

- Takes original dataset D with N training examples
- Creates M different copies $\{D_m\}_{m=1}^M$ via
 - Sampling from D with replacement
 - Each D_m has same number of examples as D
- Train **decision trees** t_1, \dots, t_M using D_1, \dots, D_m
 - suppose there are k features
 - randomly sample k' ($\leq k$) features** at each split
 - when $k' = k$ its indifferent to traditional DTs
 - usually choose $k' = \log_2 k$ (or \sqrt{k} according to Korriss)
- Average / Vote model as the final ensemble $h = \frac{1}{M} \sum_{m=1}^M h_m$

Intuition: By randomizing not just samples (step 2), but also features (step 3), we introduce even greater variety compared to *Bagging*, reducing correlation. (The entire premise of ensemble method is independence between classifiers)

Boosting

Basic Idea

$$H(x) = \sum_{t=1}^T \alpha_t h_t(x)$$

- Train T weak learners (through T iterations)
- Each weak learner:
 - only have to be slightly better than random
 - Focuses on difficult data points

Procedure

- Initialize **weights** for the N samples.
- $D_1 = \{w_{11}, \dots, w_{1N}\}$ is uniform distribution ($\forall w_{1i} = 1/N$)
- D_t means the weighted dataset at t iteration
- For $t = 1, 2, \dots, T$ do:


- Train h_t on weighted dataset D_t , and compute error of ε_t

$$\varepsilon_t = \sum_{i=1}^N P_{x \sim D_t}(h_t(x_i) \neq y_i) = \sum_{h_t(x_i) \neq y_i} w_{ti}$$

- Use h_t 's error rate ε_t to determine its weight α_t in the ensemble

$$\alpha_t = \frac{1}{2} \log \frac{1 - \varepsilon_t}{\varepsilon_t}$$

- if $\varepsilon_t > 0.5$, then $\alpha_t < 0$ and exponentially decrease.
- if $\varepsilon_t < 0.5$, then $\alpha_t > 0$ and exponentially increases.



- Use h_t 's weight α_t to determine the next weighted dataset D_{t+1}

• normalize each dimension by standard deviation

• Discrete Data: use hamming distance

Ensemble Methods

Advantages: Improved accuracy; other types of classifiers can be directly included; Easy to implement; Not too much parameter tuning

Disadvantage: Black box; Not compact representation.

Curse of Dimensionality

Definition: Prediction accuracy degrade quickly when number of attributes grow, because

- When many irrelevant attributes involved, relevant attributes are shadowed, and distance measurements are unreliable

Solution:

- Remove irrelevant attributes in pre-processing (e.g., dimensionality reduction such as PCA)

Ensemble Methods

Advantages: Improved accuracy; other types of classifiers can be directly included; Easy to implement; Not too much parameter tuning

Disadvantage: Black box; Not compact representation.

Curse of Dimensionality

Definition: Prediction accuracy degrade quickly when number of attributes grow, because

- When many irrelevant attributes involved, relevant attributes are shadowed, and distance measurements are unreliable

Solution:

- Remove irrelevant attributes in pre-processing (e.g., dimensionality reduction such as PCA)

Ensemble Methods

Advantages: Improved accuracy; other types of classifiers can be directly included; Easy to implement; Not too much parameter tuning

Disadvantage: Black box; Not compact representation.

Curse of Dimensionality

Definition: Prediction accuracy degrade quickly when number of attributes grow, because

- When many irrelevant attributes involved, relevant attributes are shadowed, and distance measurements are unreliable

Ensemble Methods

Advantages: Improved accuracy; other types of classifiers can be directly included; Easy to implement; Not too much parameter tuning

Disadvantage: Black box; Not compact representation.

Curse of Dimensionality

Definition: Prediction accuracy degrade quickly when number of attributes grow, because

- When many irrelevant attributes involved, relevant attributes are shadowed, and distance measurements are unreliable

Ensemble Methods

Advantages: Improved accuracy; other types of classifiers can be directly included; Easy to implement; Not too much parameter tuning

Disadvantage: Black box; Not compact representation.

Curse of Dimensionality

Definition: Prediction accuracy degrade quickly when number of attributes grow, because

- When many irrelevant attributes involved, relevant attributes are shadowed, and distance measurements are unreliable

Ensemble Methods

Advantages: Improved accuracy; other types of classifiers can be directly included; Easy to implement; Not too much parameter tuning

Disadvantage: Black box; Not compact representation.

Curse of Dimensionality

Definition: Prediction accuracy degrade quickly when number of attributes grow, because

- When many irrelevant attributes involved, relevant attributes are shadowed, and distance measurements are unreliable

Ensemble Methods

Advantages: Improved accuracy; other types of classifiers can be directly included; Easy to implement; Not too much parameter tuning

Disadvantage: Black box; Not compact representation.

Curse of Dimensionality

Definition: Prediction accuracy degrade quickly when number of attributes grow, because

- When many irrelevant attributes involved, relevant attributes are shadowed, and distance measurements are unreliable

Ensemble Methods

Advantages: Improved accuracy; other types of classifiers can be directly included; Easy to implement; Not too much parameter tuning

Disadvantage: Black box; Not compact representation.

Curse of Dimensionality

Definition: Prediction accuracy degrade quickly when number of attributes grow, because

- When many irrelevant attributes involved, relevant attributes are shadowed, and distance measurements are unreliable

Ensemble Methods

Advantages: Improved accuracy; other types of classifiers can be directly included; Easy to implement; Not too much parameter tuning

Disadvantage: Black box; Not compact representation.

Curse of Dimensionality

Definition: Prediction accuracy degrade quickly when number of attributes grow, because

- When many irrelevant attributes involved, relevant attributes are shadowed, and distance measurements are unreliable

Ensemble Methods

Advantages: Improved accuracy; other types of classifiers can be directly included; Easy to implement; Not too much parameter tuning

Disadvantage: Black box; Not compact representation.

Curse of Dimensionality

Definition: Prediction accuracy degrade quickly when number of attributes grow, because

- When many irrelevant attributes involved, relevant attributes are shadowed, and distance measurements are unreliable

Ensemble Methods

Advantages: Improved accuracy; other types of classifiers can be directly included; Easy to implement; Not too much parameter tuning

Disadvantage: Black box; Not compact representation.

Curse of Dimensionality

Definition: Prediction accuracy degrade quickly when number of attributes grow, because

- When many irrelevant attributes involved, relevant attributes are shadowed, and distance measurements are unreliable

Ensemble Methods

Advantages: Improved accuracy; other types of classifiers can be directly included; Easy to implement; Not too much parameter tuning

Disadvantage: Black box; Not compact representation.

Curse of Dimensionality

Definition: Prediction accuracy degrade quickly when number of attributes grow, because

- When many irrelevant attributes involved, relevant attributes are shadowed, and distance measurements are unreliable

Ensemble Methods

Advantages: Improved accuracy; other types of classifiers can be directly included; Easy to implement; Not too much parameter tuning

Disadvantage: Black box; Not compact representation.

Curse of Dimensionality

Definition: Prediction accuracy degrade quickly when number of attributes grow, because

- When many irrelevant attributes involved, relevant attributes are shadowed, and distance measurements are unreliable

Ensemble Methods

Advantages: Improved accuracy; other types of classifiers can be directly included; Easy to implement; Not too much parameter tuning

Disadvantage: Black box; Not compact representation.

Curse of Dimensionality

Definition: Prediction accuracy degrade quickly when number of attributes grow, because

- When many irrelevant attributes involved, relevant attributes are shadowed, and distance measurements are unreliable

Ensemble Methods

Advantages: Improved accuracy; other types of classifiers can be directly included; Easy to implement; Not too much parameter tuning

Disadvantage: Black box; Not compact representation.

Curse of Dimensionality

Definition: Prediction accuracy degrade quickly when number of attributes grow, because

- When many irrelevant attributes involved, relevant attributes are shadowed, and distance measurements are unreliable

Ensemble Methods

Advantages: Improved accuracy; other types of classifiers can be directly included; Easy to implement; Not too much parameter tuning

Disadvantage: Black box; Not compact representation.

Curse of Dimensionality

Definition: Prediction accuracy degrade quickly when number of attributes grow, because

- When many irrelevant attributes involved, relevant attributes are shadowed, and distance measurements are unreliable

Ensemble Methods

Advantages: Improved accuracy; other types of classifiers can be directly included; Easy to implement; Not too much parameter tuning

Disadvantage: Black box; Not compact representation.

Curse of Dimensionality

Definition: Prediction accuracy degrade quickly when number of attributes grow, because

- When many irrelevant attributes involved, relevant attributes are shadowed, and distance measurements are unreliable

Ensemble Methods

Advantages: Improved accuracy; other types of classifiers can be directly included; Easy to implement; Not too much parameter tuning

Disadvantage: Black box; Not compact representation.

Curse of Dimensionality

Definition: Prediction accuracy degrade quickly when number of attributes grow, because

- When many irrelevant attributes involved, relevant attributes are shadowed, and distance measurements are unreliable

Ensemble Methods

Advantages: Improved accuracy; other types of classifiers can be directly included; Easy to implement; Not too much parameter tuning

Disadvantage: Black box; Not compact representation.

Curse of Dimensionality

Definition: Prediction accuracy degrade quickly when number of attributes grow, because

- When many irrelevant attributes involved, relevant attributes are shadowed, and distance measurements are unreliable

Ensemble Methods

Advantages: Improved accuracy; other types of classifiers can be directly included; Easy to implement; Not too much parameter tuning

Disadvantage: Black box; Not compact representation.

Curse of Dimensionality

Definition: Prediction accuracy degrade quickly when number of attributes grow, because

- When many irrelevant attributes involved, relevant attributes are shadowed, and distance measurements are unreliable

Ensemble Methods

Advantages: Improved accuracy; other types of classifiers can be directly included; Easy to implement; Not too much parameter tuning

Disadvantage: Black box; Not compact representation.

Curse of Dimensionality

Definition: Prediction accuracy degrade quickly when number of attributes grow, because

- When many irrelevant attributes involved, relevant attributes are shadowed, and distance measurements are unreliable

Ensemble Methods

Advantages: Improved accuracy; other types of classifiers can be directly included; Easy to implement; Not too much parameter tuning

Disadvantage: Black box; Not compact representation.

Curse of Dimensionality

Definition: Prediction accuracy degrade quickly when number of attributes grow, because

- When many irrelevant attributes involved, relevant attributes are shadowed, and distance measurements are unreliable

Ensemble Methods

Advantages: Improved accuracy; other types of classifiers can be directly included; Easy to implement; Not too much parameter tuning

Disadvantage: Black box; Not compact representation.

Curse of Dimensionality

Definition: Prediction accuracy degrade quickly when number of attributes grow, because

- When many irrelevant attributes involved, relevant attributes are shadowed, and distance measurements are unreliable

Ensemble Methods

Advantages: Improved accuracy; other types of classifiers can be directly included; Easy to implement; Not too much parameter tuning

Disadvantage: Black box; Not compact representation.

Curse of Dimensionality

Definition: Prediction accuracy degrade quickly when number of attributes grow, because

- When many irrelevant attributes involved, relevant attributes are shadowed, and distance measurements are unreliable

Ensemble Methods

Advantages: Improved accuracy; other types of classifiers can be directly included; Easy to implement; Not too much parameter tuning

Disadvantage: Black box; Not compact representation.

Curse of Dimensionality

Definition: Prediction accuracy degrade quickly when number of attributes grow, because

- When many irrelevant attributes involved, relevant attributes are shadowed, and distance measurements are unreliable

Ensemble Methods

Advantages: Improved accuracy; other types of classifiers can be directly included; Easy to implement; Not too much parameter tuning

Disadvantage: Black box; Not compact representation.

Curse of Dimensionality

Definition: Prediction accuracy degrade quickly when number of attributes grow, because

- When many irrelevant attributes involved, relevant attributes are shadowed, and distance measurements are unreliable

Ensemble Methods

Advantages: Improved accuracy; other types of classifiers can be directly included; Easy to implement; Not too much parameter tuning

Disadvantage: Black box; Not compact representation.

Curse of Dimensionality

Missing Value (Worst to Best)

1. **worst:** Ignore Tuple. Unless missing label, this loss information.
2. **okay:** Manual imputation, fill constant (e.g., unknown) or mean.
3. **best:** Inference. Association Analysis to infer missingness.

Noisy Data

1. **Binning:** Discretize the variable
 - Equal-width (distance) partitioning: Divide range into N intervals of equal size. If the lowest-highest values of the attribute is A and B , interval width = $(B - A)/N$
 - Straightforward. But **sensitive to outliers or skewness**.
 - Equal-depth (frequency) partitioning: Divide range into N intervals, each containing (about) the same number of samples.
 - Good data scaling. Managing categorical variable can be tricky.
2. **Clustering:** Use distance-based clustering to find outliers
3. **Regression:** fit linear model. Data points far from fit are outliers.

Smoothing:

1. Binning: smooth by bin means or bin boundary
2. Clustering: replace noisy data with cluster centroid
3. Regression: replace noisy data with predicted value (project to line)

Stage 2: Data Transformation

1. Min-Max: $v' = \frac{v - \min}{\max - \min} (\max_{\text{new}} - \min_{\text{new}}) + \min_{\text{new}}$
• $[\min_{\text{new}}, \max_{\text{new}}]$ is the new range of the variable (e.g., $[-1, 1]$)
2. Z-score: $v' = \frac{v - \mu}{\sigma}$
3. Decimal Scaling: $v' = \frac{v}{10^j}$, where j is smallest integer such that $\max(|v'|) \leq 1$

Stage 3: Data Reduction

1. Reduction

- **Feature Selection:** choose $k < d$ important features, ignore rest.
 - ▷ Forward Search: start with \emptyset features. Each iteration, add best new feature. Essentially, **Greedy Hillclimbing**
 - ▷ Backward Search: start with all features. Each iteration, remove worst feature.
 - ▷ Decision Tree Induction: Use only the features used by DT.
- **Feature Extraction:** project d -dimensional data to $k < d$ dimension
 - ▷ Principal Component Analysis (PCA)
 - ▷ Linear Discriminant Analysis (LDA)
 - ▷ Factor Analysis
- **Numerosity Reduction:** reduce data vol. w/ compact representation
 - ▷ Parametric: Fit data w/ model. Store just parameters, discard data.
 - ▷ Non-parametric: histograms, clustering, (stratified) sampling.

2. Discretization

- Divide range of continuous attribute into intervals (e.g., by binning)
 - ▷ Why?
 - Some algo only accept categorical attributes
 - Reduce data size. Prepare for further analysis.
- Concept hierarchy: replace low-level concept w/ higher level concepts (e.g., “age = 25” → “young”)

Data Warehouse

	OLTP	OLAP
users	clerk, IT professional	knowledge worker
function	day to day operations	decision support
DB design	application-oriented	subject-oriented
data	current, up-to-date detailed, flat relational isolated	historical, summarized, multidimensional integrated, consolidated
usage	repetitive	ad-hoc
access	read/write index/hash on prim. key	lots of scans
unit of work	short, simple transaction	complex query
# records accessed	tems	millions
#users	thousands	hundreds
DB size	100MB-GB	100GB-TB
metric	transaction throughput	query throughput, response



Core Concepts

Definition: A *subject-oriented, integrated, time-variant, nonvolatile* collection of data to support management decision making.

OLTP vs OLAP:

- **OLTP (Online Transaction Processing):**
 - ▷ Access: Read/Write, atomic transactions.
 - ▷ Focus: Concurrency (ACID).
 - ▷ Data: Current, detailed, relational (3NF). Users: Clerks, DBAs.
- **OLAP (Online Analytical Processing):**
 - ▷ Access: Read-only (mostly), complex aggregate queries.
 - ▷ Focus: Throughput.
 - ▷ Data: Historical, consolidated, multi-dim.
 - ▷ Users: Managers, Analysts.

Multidimensional Data Model



- **MOLAP:** Multi-dim Array (Direct cube). Fast, sparse, limited scale.
- **HOLAP:** Hybrid (Detail in ROLAP, Aggregates in MOLAP).

Web Mining

Input: Unstructured/Semi-structured Data. **Goal:** Discover patterns.

1. Information Retrieval (Text DB)

Data Abstraction (Feature Eng):

- **Bag of Words:** Document d as set of terms. Ignore order.
- **Vector Space Model:** $d \in \mathbb{R}^{|V|}$ (High-dim, sparse).
 - ▷ Value = Term Frequency (TF) or TF-IDF.
 - ▷ **Preprocessing:**
 - Stop list (remove ‘the’),
 - Stemming (‘mining’ → ‘mine’).

Similarity (Metric):

- Euclidean fails in high-dim (curse of dimensionality) & length bias.
- **Cosine Similarity:** Measure angle (direction), not magnitude.

$$\text{Sim}(d_1, d_2) = \frac{d_1 \cdot d_2}{\|d_1\| \times \|d_2\|}$$

Evaluation (Set Theory):

- **Precision:** $\frac{|\{\text{Relevant} \cap \{\text{Retrieved}\}|}}{|\{\text{Retrieved}\}|}$ (Quality of result).
- **Recall:** $\frac{|\{\text{Relevant} \cap \{\text{Retrieved}\}|}}{|\{\text{Relevant}\}|}$ (Coverage of truth).

Challenges (Why Keyword Match Fails):

- **Synonymy:** Different words, same meaning (e.g., “ML” vs “Learning”). → Low Recall.
- **Polysemy:** Same word, different meanings (e.g., “Apple”). → Low Precision.

2. Taxonomy (The 3 Pillars)

Content Mining (NLP):

- **Input:** HTML Text. **Task:** Classification/Clustering.
- **Logic:** Text Mining + HTML Tags weight.

Structure Mining (Graph):

- **Input:** Links (Edges). **Task:** Find Authority/Hubs.
- **Logic:** “Link as Vote”. Graph Theory.

Usage Mining (Logs/BD):

- **Input:** Server Logs (IP, Time, URL). **Task:** User Profiling.
- **Logic:** ETL → Sessionize → Pattern Mining.

3. Core Algorithms

PageRank (Google):

- **Intuition:** Link=endorsement. Importance via recursive voting.
- **Logic:** Random Surfer Model / Markov Chain Stationary Dist.

$$\text{PR}(A) = (1 - d) + d \sum_{T_i \rightarrow A} \frac{\text{PR}(T_i)}{C(T_i)}$$

HITS (Hubs & Authorities):

- **Hub:** Guide page pointing to many authorities.
- **Authority:** Content page pointed to by many hubs.
- Mutually reinforcing: Good Hub → points to good Auth.

Web Usage Process:

1. **Clean:** Remove crawlers, graphics requests.
2. **Identify:** Split IP streams into User Sessions.
3. **Pattern:** Apriori / Sequential Mining on sessions.

Exam Questions

Network → Matrix → DBSCAN

1. Analysis and Parameter Proposal

The dissimilarity matrix (Fig. 4) represents **structural equivalence** (Hamming distance of adjacency vectors) rather than path length. Notably, hubs D and E have high dissimilarity to others ($dge3$) and to each other ($d = 6$).

To distinguish tight structural groups from loose connections:

- **Radius (ϵ) = 2:** A threshold of $\epsilon \geq 3$ would merge the entire graph into a single cluster (since $d(A, D) = 3, d(A, E) = 3$). $\epsilon = 2$ exploits the natural gap in the data to separate strong ties.
- **MinPts = 2:** Chosen to allow the detection of the smallest possible social units (dyads) visible in the graph (e.g., A-F).

2. Execution (Neighborhood Analysis)

We calculate the ϵ -neighborhood $N_{(\epsilon)}(x) = \{ym \mid d(x, y) \leq \epsilon\}$ for each node.

Node	Neighborhood (dle2)	Status (MinPts=2)
A	{A, F} (Size: 2)	Core Point
F	{F, A} (Size: 2)	Core Point
B	{B, C} (Size: 2)	Core Point
C	{C, B} (Size: 2)	Core Point
D	{D} (Size: 1, min dist is 3)	Noise
E	{E} (Size: 1, min dist is 3)	Noise

3. Final Clustering Result

DBSCAN expands clusters from the core points by merging overlapping neighborhoods.

Implementation & Optimization

1. Partial Materialization (Cube Computation):

- **Problem:** 2^N cuboids. Full materialization too big. No mat. too slow.
- **Solution:** Greedy Algo. Select views with max benefit/space ratio.
- **Iceberg Cube:** Only store cells w/ value $>$ threshold (sparsity).

2. Indexing Techniques:

- **Bitmap Index:** For low-cardinality columns (Gender). fast.
- **Join Index:** Map Join(Fact.key, Dim.key) to avoid runtime joins.

3. Architecture:

- **ROLAP:** Relational (Star schema). Scalable, slower.

• Cluster 1: {A, F} (High structural similarity)

• Cluster 2: {B, C} (High structural similarity)

• Outliers: D and E. Despite being central graph nodes, they are **structurally unique** (dissimilar neighbors) and fail to meet the density criteria.

Datawarehouse

1. Missing Value Imputation

The missing values are filled based on the class defined by attributes "Fever" and "Disease".

1.1. Imputing Age (Patient 9910115)

- Class Definition: Fever = "No", Disease = "Yes".
- Matching Records:

• Patient 9303034: Age 55

• Patient 9910111: Age 46

• Calculation (Mean): $\frac{55+46}{2} = 50.5$

1.2. Imputing Sex (Patient 9576737)

- Class Definition: Fever = "Yes", Disease = "No".
- Matching Records:

• Patient 9100123: Male

• Patient 9576732: Male

• Calculation (Mode): Male

Answer:

- Missing Age: **50.5**
- Missing Sex: **Male**

2. Data Smoothing by Bin Means

Raw Data (Age): {65, 55, 12, 35, 46, 16, 105, 28}

Step 1: Sort the Data

Step 2: Partition into 4 Equi-depth Bins Total items = 8. Items per bin = $\frac{8}{4} = 2$.

• Bin 1: {12, 16}

• Bin 2: {28, 35}

• Bin 3: {46, 55}

• Bin 4: {65, 105}

Step 3: Calculate Bin Means and Smooth

• Bin 1 Mean: $\frac{12+16}{2} = 14$

• Bin 2 Mean: $\frac{28+35}{2} = 31.5$

• Bin 3 Mean: $\frac{46+55}{2} = 50.5$

• Bin 4 Mean: $\frac{65+105}{2} = 85$

Smoothed Values List: Bin 1: 14, 14 Bin 2: 31.5, 31.5 Bin 3: 50.5,

50.5 Bin 4: 85, 85

3. 3D Data Cube Design

Context: Hospital Revenue Analysis. Star Schema Structure:

1. Dimension Tables:

- Time Dimension: Attributes include Time_Key, Day, Month, Quarter, Year. Provides temporal granularity.
- Location Dimension: Attributes include Location_Key, Department_Name, Hospital_Branch, City, Region. Provides geographical/organizational hierarchy.
- Treatment Dimension: Attributes include Treatment_Key, Procedure_Code (e.g., CPT), Category (e.g., Surgery, Pathology), Risk_Level. Describes the service provided.

2. Fact Table:

- Revenue_Fact: Contains foreign keys linking to the three dimensions (Time_Key, Location_Key, Treatment_Key) and numerical measures such as Total_Revenue (\$) and Patient_Count.

4. OLAP Operations

Based on the context in (c), assume the current query view is: "Total Revenue by Month and Hospital Branch."

1. Roll Up (Aggregation/Generalization)

- Operation: Climb up the concept hierarchy on the Time dimension.
- Action: Change grouping from Month to Quarter or Year.
- Result: The data becomes less detailed, showing total revenue per Quarter/Year, summarizing the monthly variations.

2. Drill Down (Detailed View/Specialization)

- Operation: Step down the concept hierarchy on the Location dimension.
- Action: Expand Hospital_Branch to Department_Name.
- Result: The data becomes more detailed, breaking down the branch's revenue to show specifically which departments (e.g., ER, Oncology, Pediatrics) generated the funds.

Sequence Mining

Section B [80%]: Long Questions

B1. [20 marks]

Given the following tourist data records... (Refer to original image for full dataset) Mappings: Hamburg (H), Toronto (T), Osaka (O), Beijing (B), London (L), Vancouver (V), Sydney (S).

Tourist ID	Sequence
30001	V → T → O → B → L
30002	T → B → O → T → H
30003	O → T → B → O
30004	B → H

30005	B → S → O → H → O → H
-------	-----------------------

$\langle B, O, H, O, H \rangle$

d) Compute the support, confidence and interest of the following sequential association rule.

Beijing Osaka ⇒ Hamburg

(6 marks)

Answer:

We interpret the rule as: Given a sequence contains Beijing (B) followed eventually by Osaka (O), does it also contain Hamburg (H) eventually after O?

• Rule: $\langle B, O \rangle \Rightarrow \langle H \rangle$

• Composite Sequence: $\langle B, O, H \rangle$

1. Calculate Counts (N=5):

• Antecedent Support Count ($\langle B, O \rangle$):

- 30001: $\langle V, T, O, B, L \rangle$ (Order is O → B, not B → O) → No.
- 3002: $\langle T, B, O, T, H \rangle$ → Yes (1).
- 3003: $\langle O, T, B, O \rangle$ → Yes (2).
- 3004: $\langle B, H \rangle$ (No O) → No.
- 3005: $\langle B, S, O, H, O, H \rangle$ → Yes (3).

• Rule Support Count ($\langle B, O \rangle \rightarrow \langle H \rangle$):

- 30002: $\langle \dots B, O, \dots H \rangle$ → Yes (1).
- 3003: Contains B, O, but no H. → No.
- 3005: $\langle \dots B, O, \dots H \rangle$ → Yes (2).

• Count(Rule) = 2.

• Consequent Support Count (H occurs in sequence):

- Appears in 30002, 30004, 30005.

• Count(H) = 3.

2. Compute Metrics:

• Support (s):

$$s = \frac{\text{Count}(\langle B, O, H \rangle)}{N} = \frac{2}{5} = 0.4 \text{ or } 40\%$$

• Confidence (c):

$$c = \frac{\text{Count}(\langle B, O, H \rangle)}{\text{Count}(\langle B, O \rangle)} = \frac{2}{3} \approx 0.67 \text{ or } 66.7\%$$

• Interest (Lift):

$$\text{Interest} = \frac{\text{Confidence}}{P(\text{Consequent})} = \frac{2/3}{3/5} = \frac{2}{3} \times \frac{5}{3} = \frac{10}{9} \approx 1.11$$

Naive Bayes

Outlook	Temp	Humid	Wind	Play?
Sunny	Hot	High	Weak	No
Sunny	Hot	High	Strong	No
Overcast	Hot	High	Weak	Yes
Rain	Mild	High	Weak	Yes
Rain	Cool	Normal	Weak	Yes
Rain	Cool	Normal	Strong	No
Overcast	Cool	Normal	Strong	Yes
Sunny	Mild	High	Weak	No
Sunny	Cool	Normal	Weak	Yes
Rain	Mild	Normal	Weak	Yes
Sunny	Mild	Normal	Strong	Yes
Overcast	Mild	High	Strong	Yes
Overcast	Hot	Normal	Weak	Yes
Rain	Mild	High	Strong	No

Query: X = (Sunny, Cool, High, Strong) – Play?

Step 1: Prior Probabilities

$$P(\text{Yes}) = \frac{9}{14}, P(\text{No}) = \frac{5}{14}$$

Step 2: Likelihoods (from training counts)

Feature	P(feat Yes)	P(feat No)
Sunny	2/9	3/5
Cool	3/9	1/5
High	3/9	4/5
Strong	3/9	3/5

Step 3: Apply Naive Bayes

$$P(X|\text{Yes}) = \frac{2}{9} \cdot \frac{1}{9} \cdot \frac{3}{9} \cdot \frac{3}{9} = \frac{54}{6561} \approx 0.0082$$

$$P(X|\text{No}) = \frac{3}{5} \cdot \frac{1}{5} \cdot \frac{4}{5} \cdot \frac{3}{5} = \frac{36}{625} \approx 0.0576$$

Unnormalized posteriors:

$$P(\text{Yes}) \cdot P(X|\text{Yes}) = \frac{9}{14} \cdot \frac{54}{6561} \approx 0.0053$$

$$P(\text{No}) \cdot P(X|\text{No}) = \frac{5}{14} \cdot \frac{36}{625} \approx 0.0206$$

Step 4: Normalize & Classify

$$P(\text{Yes}|X) = \frac{0.0053}{0.0053 + 0.0206} \approx 20.4\%$$

$$P(\text{No}|X) = \frac{0.0206}{0.0053 + 0.0206} \approx 79.6\%$$

Prediction: No (Don't Play Tennis)

Laplace Smoothing (for zero counts)

$$P(x_i | C) = \frac{N_{x_i, C} + \alpha}{N_C + \alpha \cdot k}$$

where $\alpha = 1$ (typically), k = number of feature values.

Transformed Sequence: