

COMP4434 Group Project

Group 5 Project Presentation

WANG Yuqi - [REDACTED]

YANG Xikun - [REDACTED]

CUI Mingyue - [REDACTED]

Contents



1	Selected Paper	3
1.a	Basic Information	4
1.b	Challenge	5
1.c	Solution: World Model	6
1.d	Dreamer V3	7
1.e	Datasets (Environments)	9
2	Designed Model	10
2.a	Hypothesized Limitation	11
2.b	Proposed Method - Pipeline Overview	14
2.c	Proposed Method - World Model	15
2.d	Proposed Method - RL Agent	16

1 Selected Paper

Mastering Diverse Domains through World Models



Journal: Nature 2025
Author: Danijar Hafner, Jurgis Pasukonis, Jimmy Ba, Timothy Lillicrap
Scope: Sample-efficient RL via self-supervised learning (World Models)

- **Reinforcement Learning is “BAD”**
 - Lacked generality and robustness
 - Sample Inefficiency - tremendous trial & error
 - Sparse Reward (Credit Assignment Problem)
- **Research Problem**
 - How can we design a sample efficient RL algorithm?
 - How can we overcome sparse reward + high dimension problem

Learning Abstract World Models

Learning a description of how the world works

- Predict future outcomes to enable **planning**
- Learn **compact representation** of images
- Form **Markov states** in POMDPs

Application

- **Exploration** by planning from expected surprise
- **improved generalization** from fixed datasets
- Solve many **new tasks** zero-shot

Main Challenge

- Learning **accurate world models** is challenging!!!
- **DreamerV3**: one of the SOTA world model based RL techniques




Figure 1: World Model learning

1. World Model

- Encoder: $x_t \rightarrow z_t$ (**categorical discrete latent**)
- RSSM: $(h_{t-1}, z_{t-1}, a_{t-1}) \rightarrow h_t$
- Decoder: $z_t \rightarrow \hat{x}_t$ (reconstruction)
- **Purpose:** Generate massive number of imagined rollouts in latent space




Figure 2: Actor-Critic

2. Reinforcement Learning

- Actor-Critic Learning based on the imagined rollouts generated by the world model.


3. Various other tricks

- Symlog, Discrete World Models, KL Balancing etc.


Datasets (Environments)



- **DeepMind Control Suite**
 - ▶ 18 Control tasks based on proprioceptive vector input
 - ▶ 20 Continuous tasks based on RGB image input
- **Atari** - 26 Atari Games
- **ProcGen** - 16 procedurally generated games
- **DMLab** - 30 3D environment tasks
- **Minecraft** - Open-world 3D sandbox game



(a) Control Suite



(b) Atari



(c) ProcGen



(d) DMLab



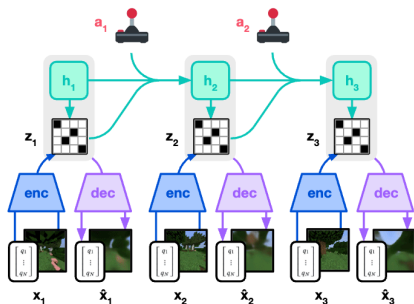
(e) Minecraft

For our project: 18 proprio-control tasks from Deepmind Control Suite


- Reason: Least computationally demanding

2 Designed Model

DreamerV3 Approach:



(a) World Model Learning



(b) Actor Critic Learning

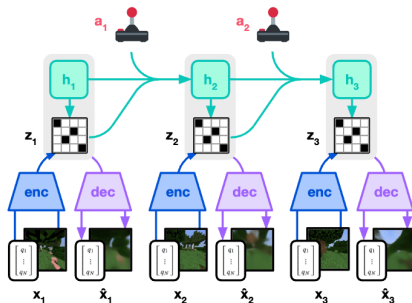
RECALL: World Model Definition

- Learning a description of how the world works
- Learn **compact representation** of images
- Form **Markov states** in POMDPs


Hypothesized Limitation (ii)



DreamerV3 Approach:



(a) World Model Learning



(b) Actor Critic Learning

Hypothesized Issues:

- Autoencoder inevitably encodes irrelevant details to RL learning
- Pixel uncertainty → necessitate **engineering tricks** for uncertainty modeling

Our Assumptions

- If predict at a sufficiently abstract and high-level representation space
- **No need for uncertainty modeling!**

An Extreme Example:

- You can deterministically predict planetary orbit with 6 variables
 - 3 variables for 3D position
 - 3 variables for velocity vector
- But if predicting planetary orbit at **atomic level**:
 - Tremendous uncertainty!
 - Error accumulation after each particle simulation

Proposed Method - Pipeline Overview



Producer:

```
while is_running:
    episode = []
    while not done:
        latent = world_model.encode(obs)
        action = agent.act(latent)
        obs, reward, done = env.step(action)
        episode.append((latent, action, reward))

    replay_buffer.add(episode)
```

Consumer:

```
while is_running:
    episode = replay_buffer.sample()
    world_model.train(episode)
    agent.train(episode)
```

Proposed Method - World Model



Components

- Encoder: $x_t \rightarrow z_t$
- Predictor: $z_t, a_t \rightarrow z_{t+1}, r_{t+1}$

Given Variables

- Observed trajectory: $[x_0, x_1, \dots, x_t]$
- Action sequences: $[a_0, a_1, \dots, a_t]$

Contrastive Training Procedure (`world_model.train()`)

- **Encoder**: encodes trajectory into latents: $[x_0, x_1, \dots, x_t] \rightarrow [z_0, z_1, \dots, z_t]$
- **Predictor**: predicts $[\hat{z}_1, \hat{z}_2, \dots, \hat{z}_t]$ from $[z_0] + [a_0, a_1, \dots, a_{t-1}]$
- Compute $\cos(z, \hat{z})$ cosine similarity matrix:

$$\begin{pmatrix} \cos(z_1, \hat{z}_1) & \dots & \cos(z_1, \hat{z}_t) \\ \vdots & \ddots & \vdots \\ \cos(z_t, \hat{z}_1) & \dots & \cos(z_t, \hat{z}_t) \end{pmatrix}$$

- `F.cross_entropy` maximize value on the diagonal, and minimize everything else

Intuition

- This **avoids reconstruction** in autoencoders
- Speeds up training and prevent encoding irrelevant visual details

Proposed Method - RL Agent



Algorithm:

- Proximal Policy Optimization (PPO)

During Training

- Data: z_t, a_t, z_{t+1}, r_t (state, action, next state, reward)
- Input: z_t (current observation/state)
- Output: $a_t \sim \pi(\cdot | z_t), V(z_t)$ (sampled action & state value estimate)
- Training objective:

$$L_{\text{total}} = \text{clamp}(L_{\text{CLIP}}, 1 - \varepsilon, 1 + \varepsilon) + c_1 L_{\text{VF}} - c_2 S[\pi(\cdot | s_t)]$$

Where:

- L_{CLIP} : Clipped policy gradient
- L_{VF} : Value function MSE loss
- S : Entropy regularization

During Inference $z_t \rightarrow a_t \sim \pi(\cdot | z_t)$ (single-step action selection)